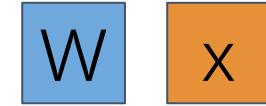


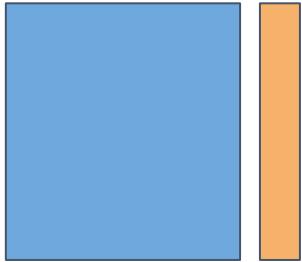
Tensor field networks: Rotation- and translation-equivariant neural networks for 3D point clouds

Nathaniel Thomas, Tess Smidt, Steven Kearnes .etc.

Neural networks are specially designed for different data types.
Assumptions about the data type are built into how the network operates.

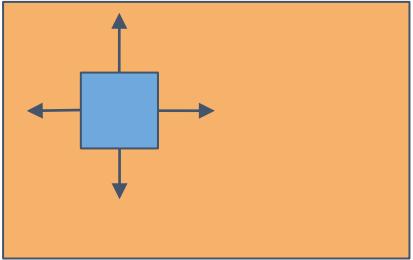


Arrays \diamond *Dense NN*



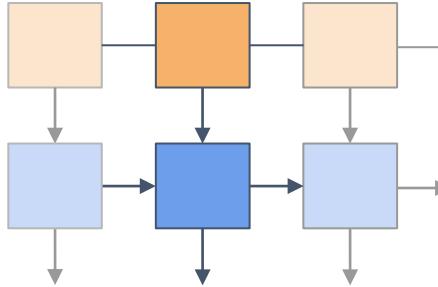
Components are independent.

2D images
 \diamond *Convolutional NN*



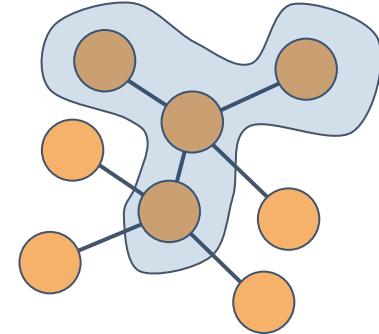
The same features can be found anywhere in an image. Locality.

Text \diamond *Recurrent NN*



Sequential data. Next input/output depends on input/output that has come before.

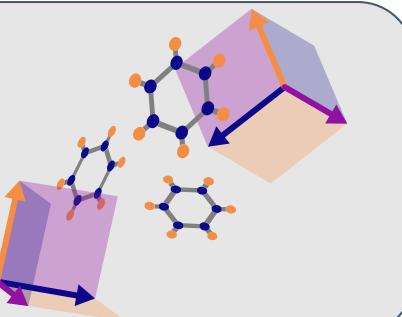
Graph \diamond *Graph (Conv.) NN*



Topological data. Nodes have features and network passes messages between nodes connected via edges.

3D physical data
 \diamond *Euclidean NN*

Data in 3D Euclidean space.
Freedom to choose coordinate system.



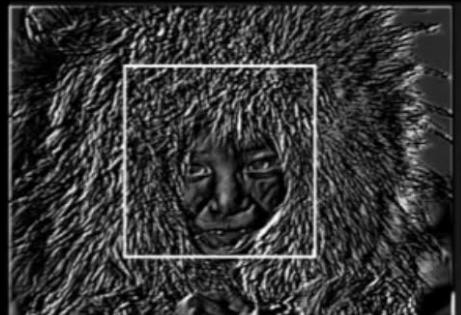
蜜蜂剪辑

Existing CNNs: Translation Equivariance

Input



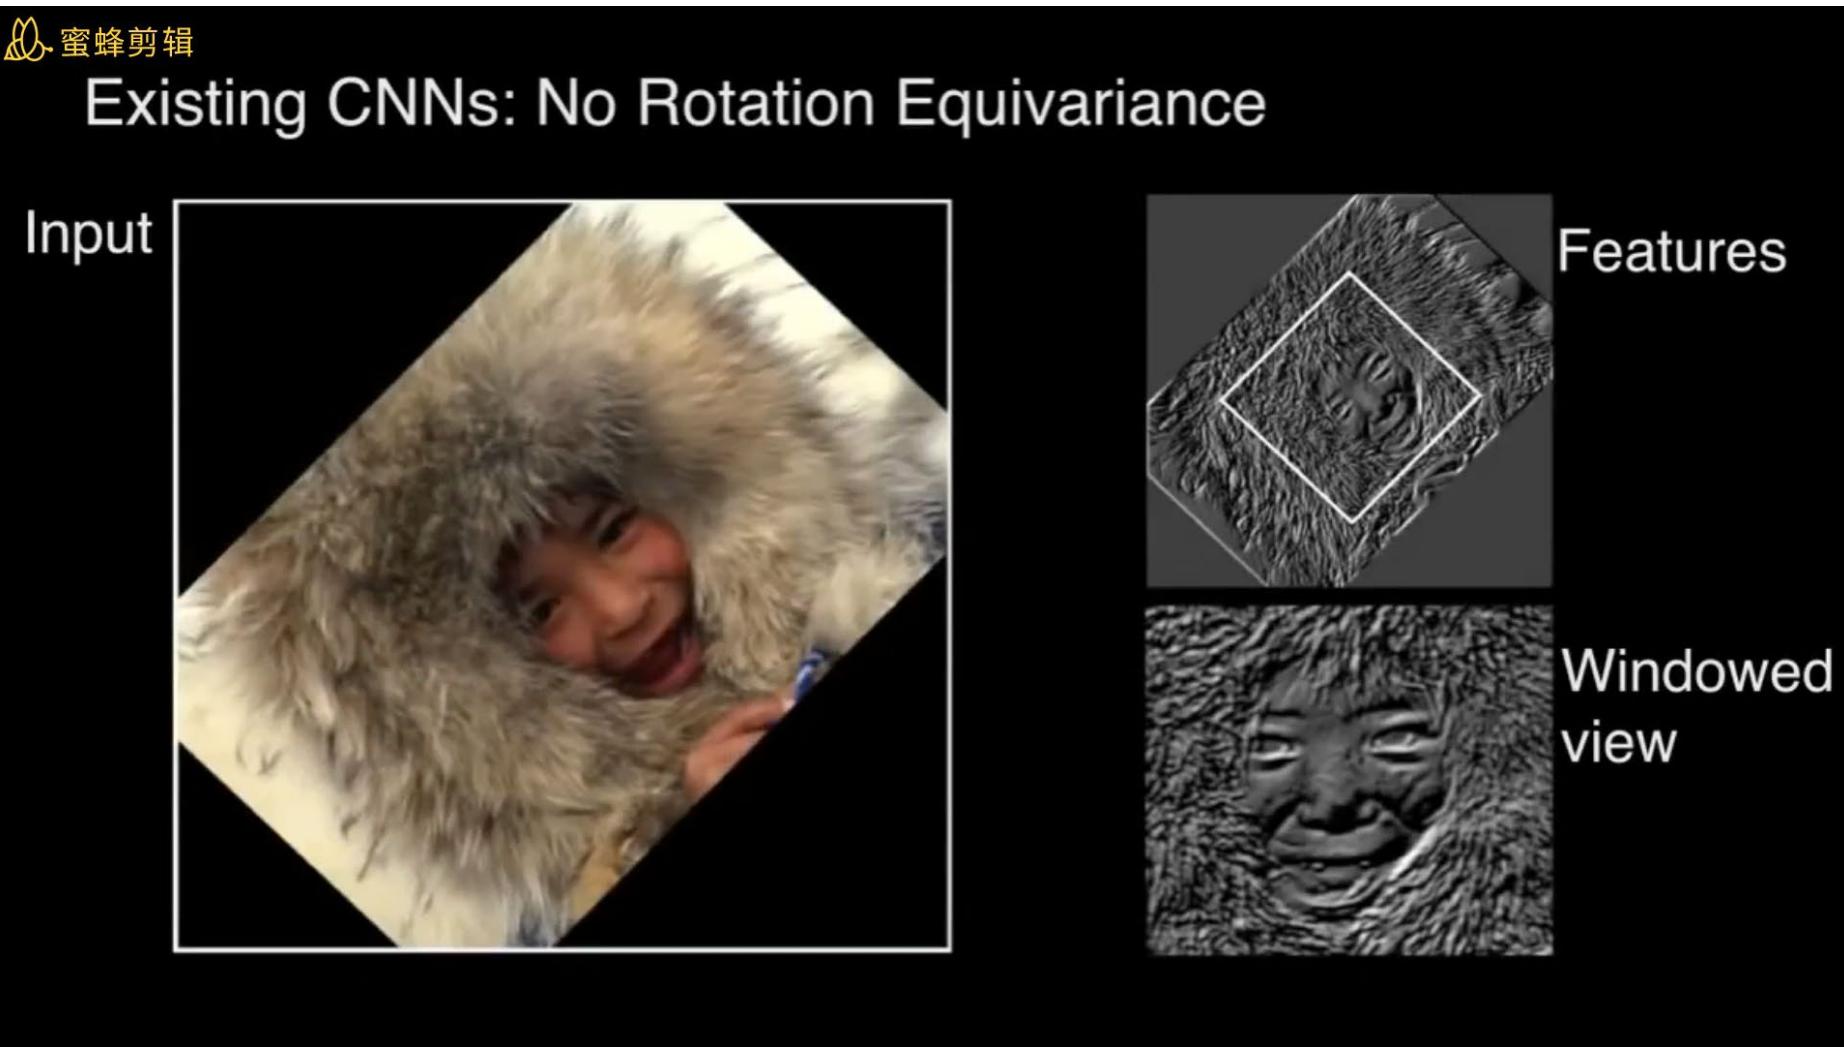
Features



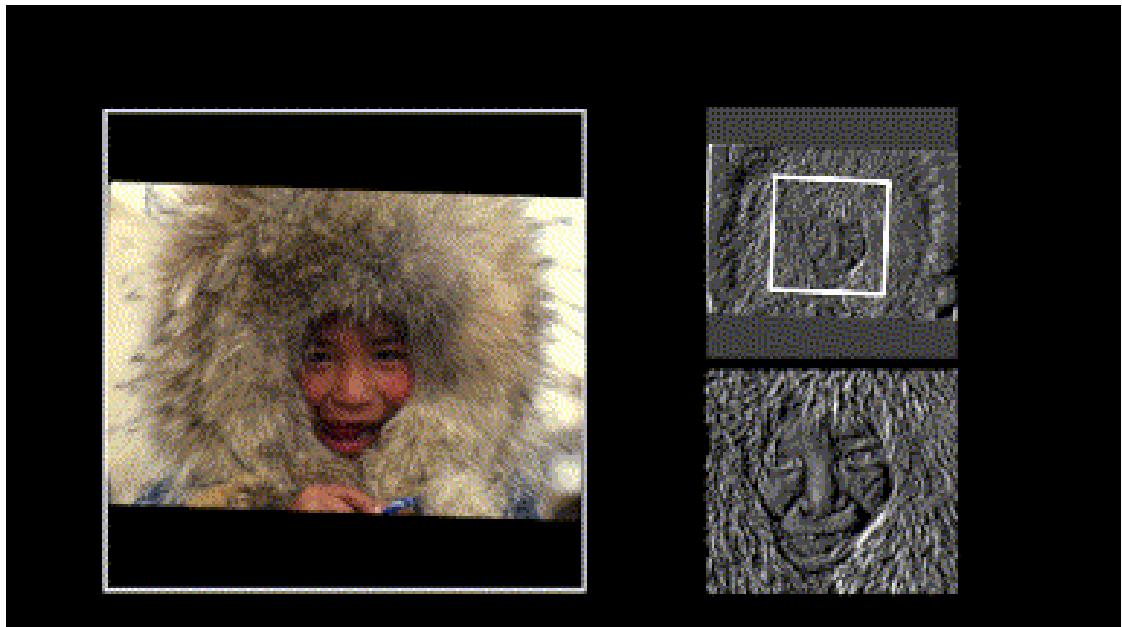
Windowed
view



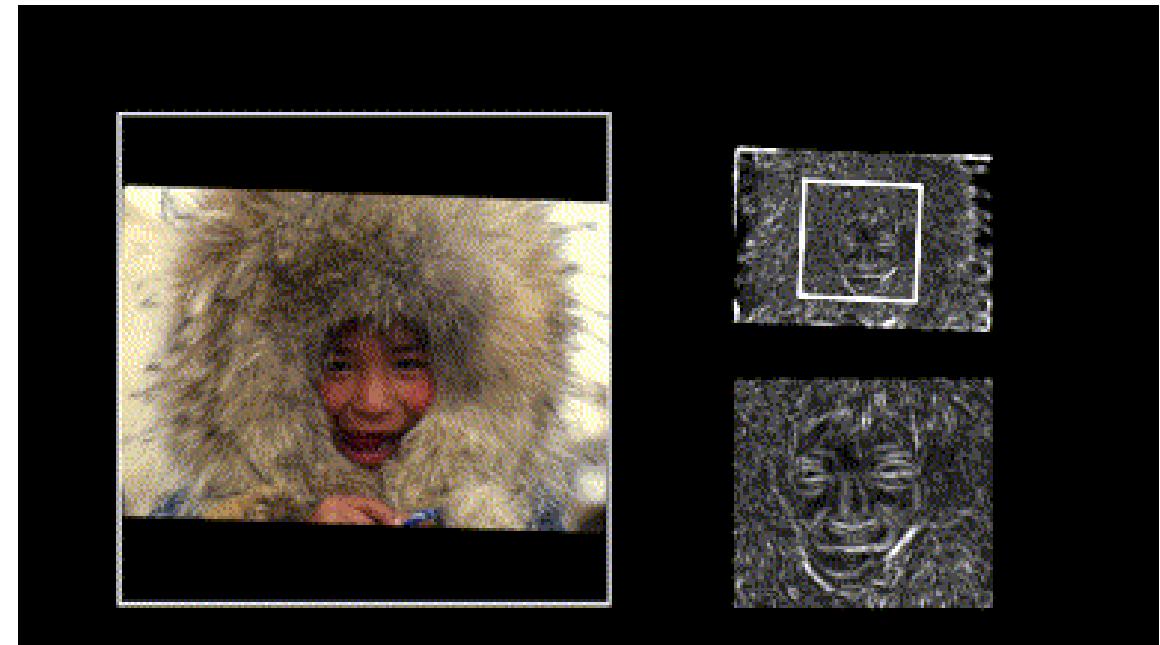
CNN-LACK OF KNOWN ROTATION EQUIVARIANCE:



CNN vs Hnet

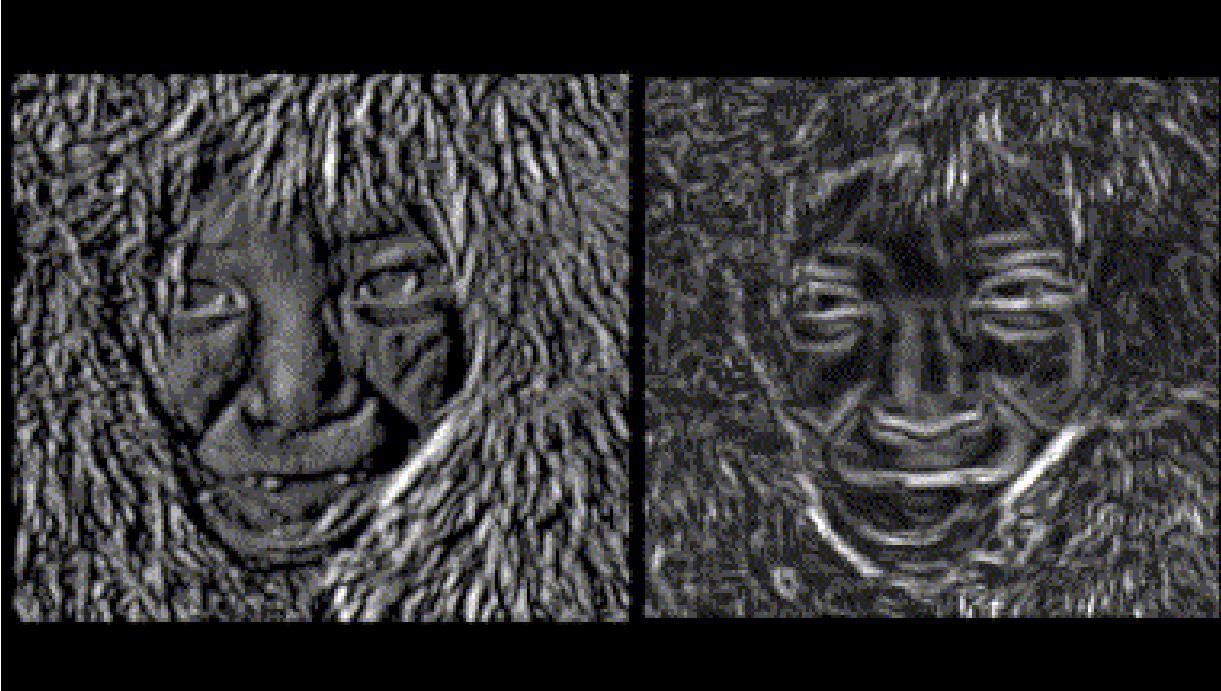


Regular CNN



Regular HNet

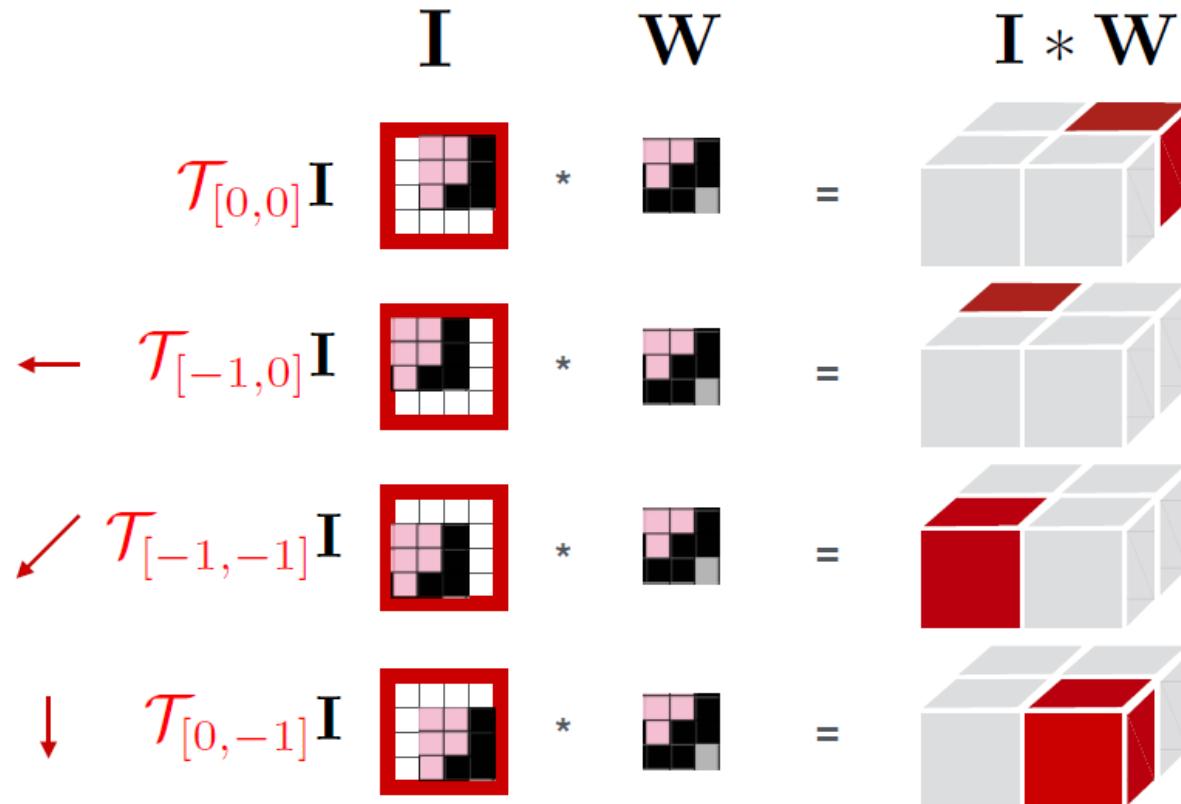
Equivalent



CONVOLUTIONS ON TRANSLATION GROUP

$$[\mathbf{I} * \mathbf{W}](\mathbf{y}) = \sum_{\mathbf{x} \in \mathbb{Z}^2} \mathbf{I}(\mathbf{x}) \mathbf{W}(\mathbf{x} - \mathbf{y})$$

Translating kernel

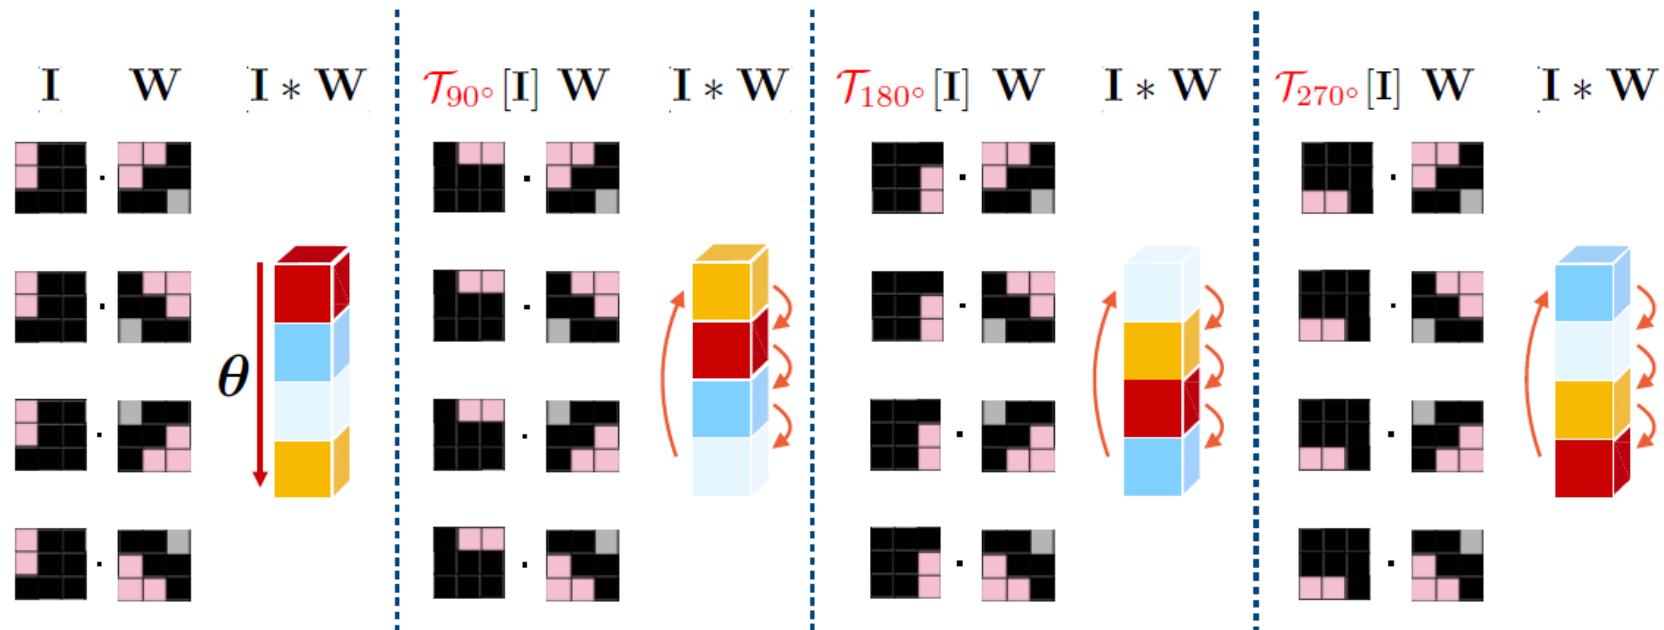


CONVOLUTIONS ON ROTATION GROUP angle

Group convolution **Notice index of convolution** Ref. $[I * W](y) = \sum_{x \in \mathbb{Z}^2} I(x)W(x - y)$

$$[I * W](\theta) = \sum_{x \in \mathbb{Z}^2} I(x)W(R_\theta^{-1}x)$$

Rotating kernel



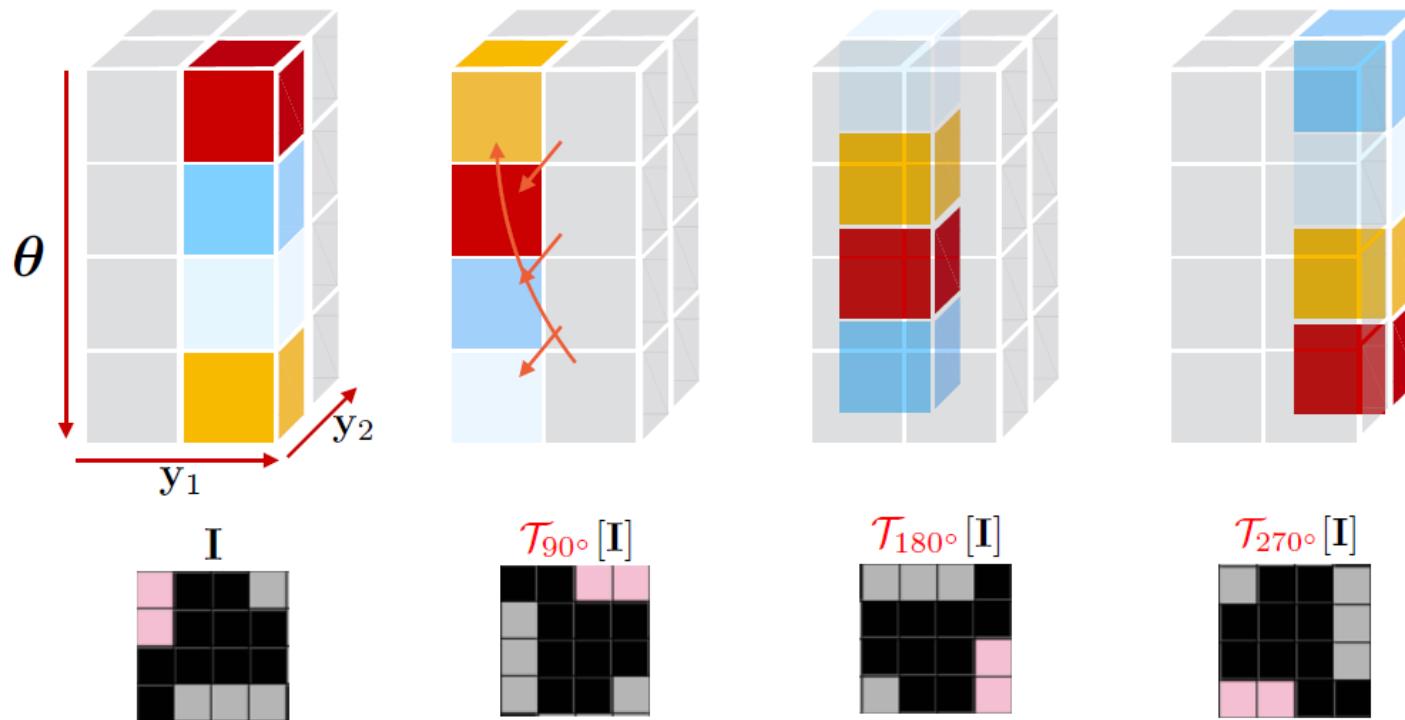
@ hidden layer $[I * W](\psi) = \sum_{\theta \in \Theta} I(\theta)W(\theta - \psi)$

CONVOLUTIONS ON ROTO-TRANSLATION

Group convolution

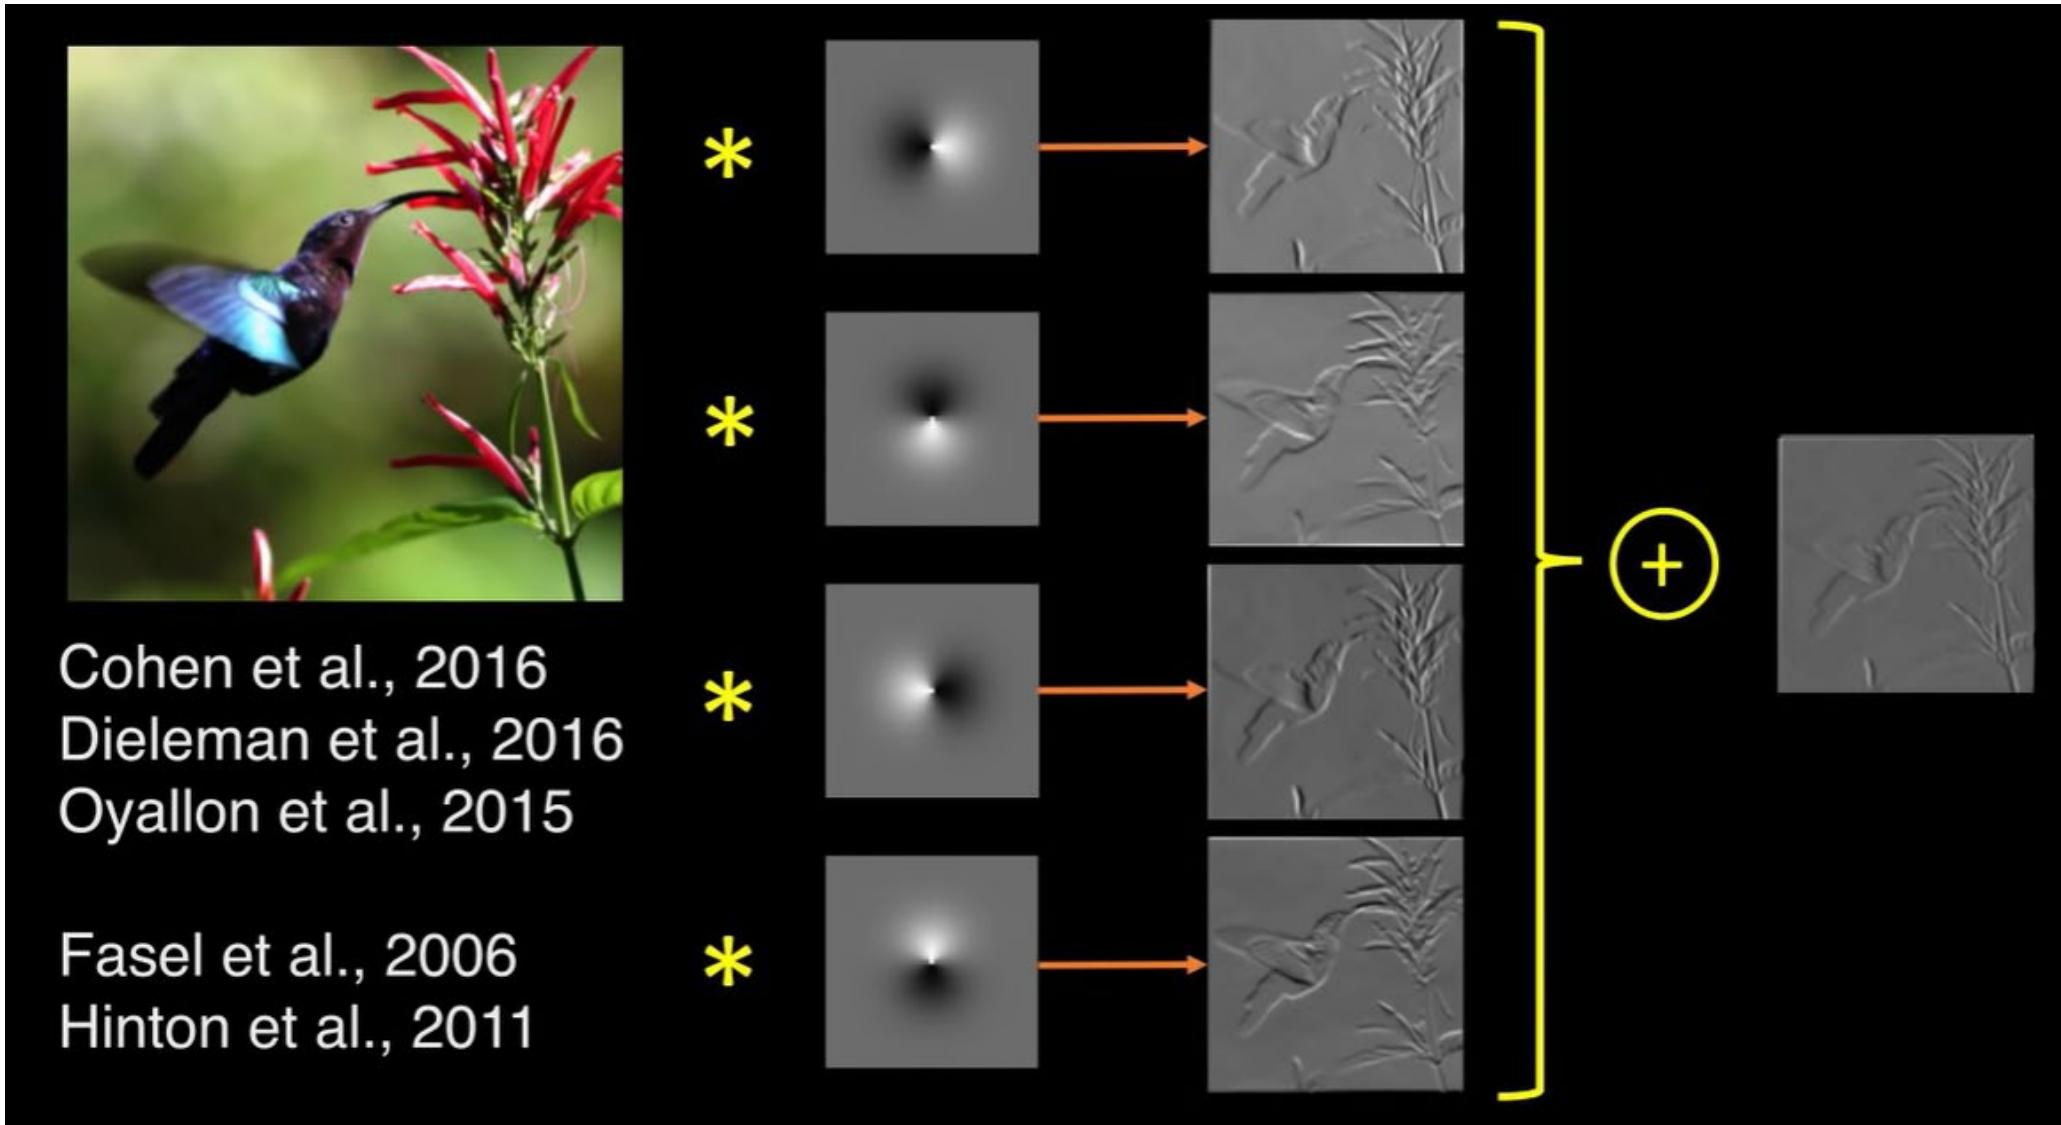
$$[\mathbf{I} * \mathbf{W}](\theta, \mathbf{y}) = \sum_{\mathbf{x} \in \mathbb{Z}^2} \mathbf{I}(\mathbf{x}) \mathbf{W}(\mathbf{R}_\theta^{-1} \mathbf{x} - \mathbf{y})$$

Roto-translating kernel

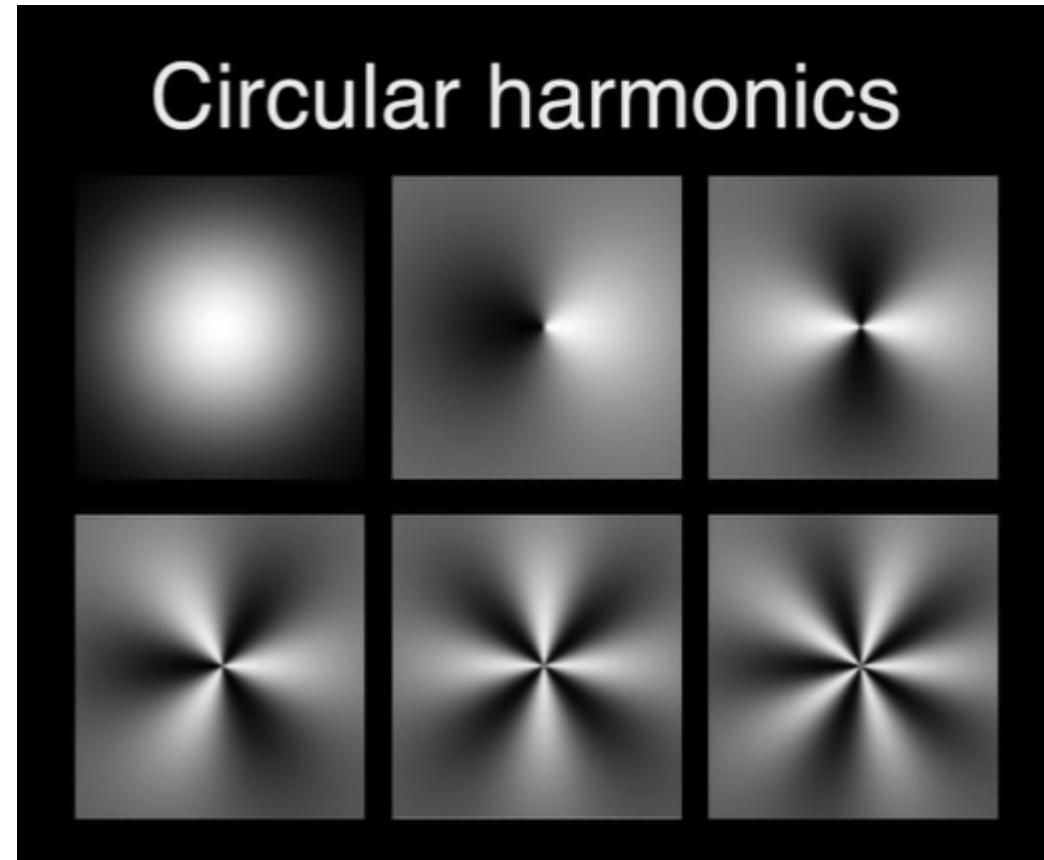
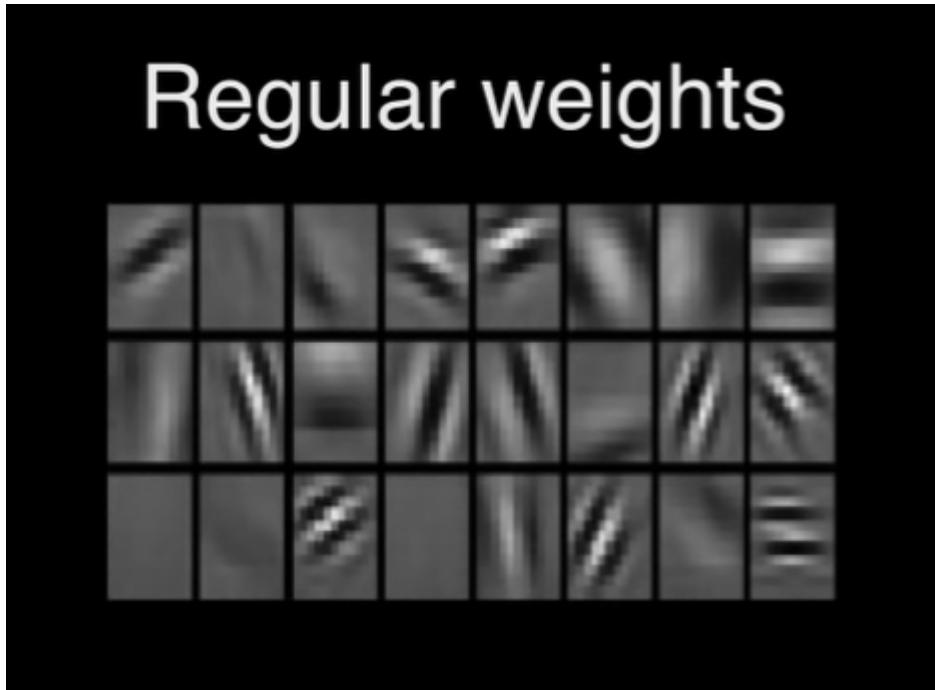


@ hidden layer $[\mathbf{I} * \mathbf{W}](\psi, \mathbf{z}) = \sum_{(\theta, \mathbf{z}) \in (\Theta \times \mathbb{Z}^2)} \mathbf{I}(\theta, \mathbf{y}) \mathbf{W}(\mathbf{R}_\psi^{-1} \mathbf{y} - \mathbf{z}, \theta - \psi)$

CONVOLUTIONS ON ROTO-TRANSLATION-Harmonic Networks-discrete



CONVOLUTIONS ON ROTO-TRANSLATION-Harmonic Networks



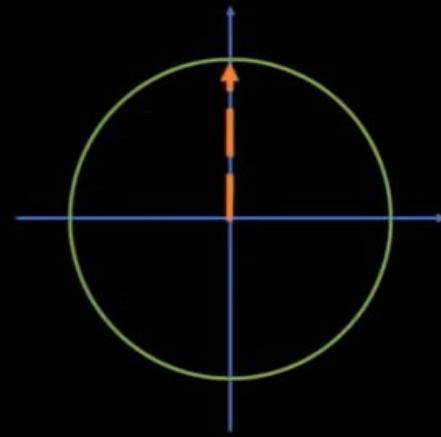
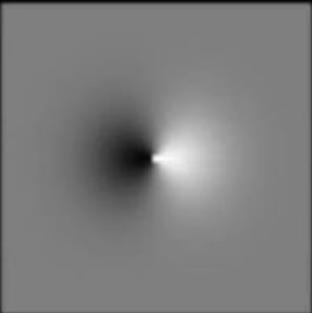
continuous

CONVOLUTIONS ON ROTO-TRANSLATION-Harmonic Networks

Harmonic Networks



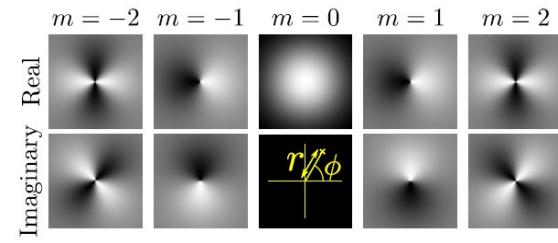
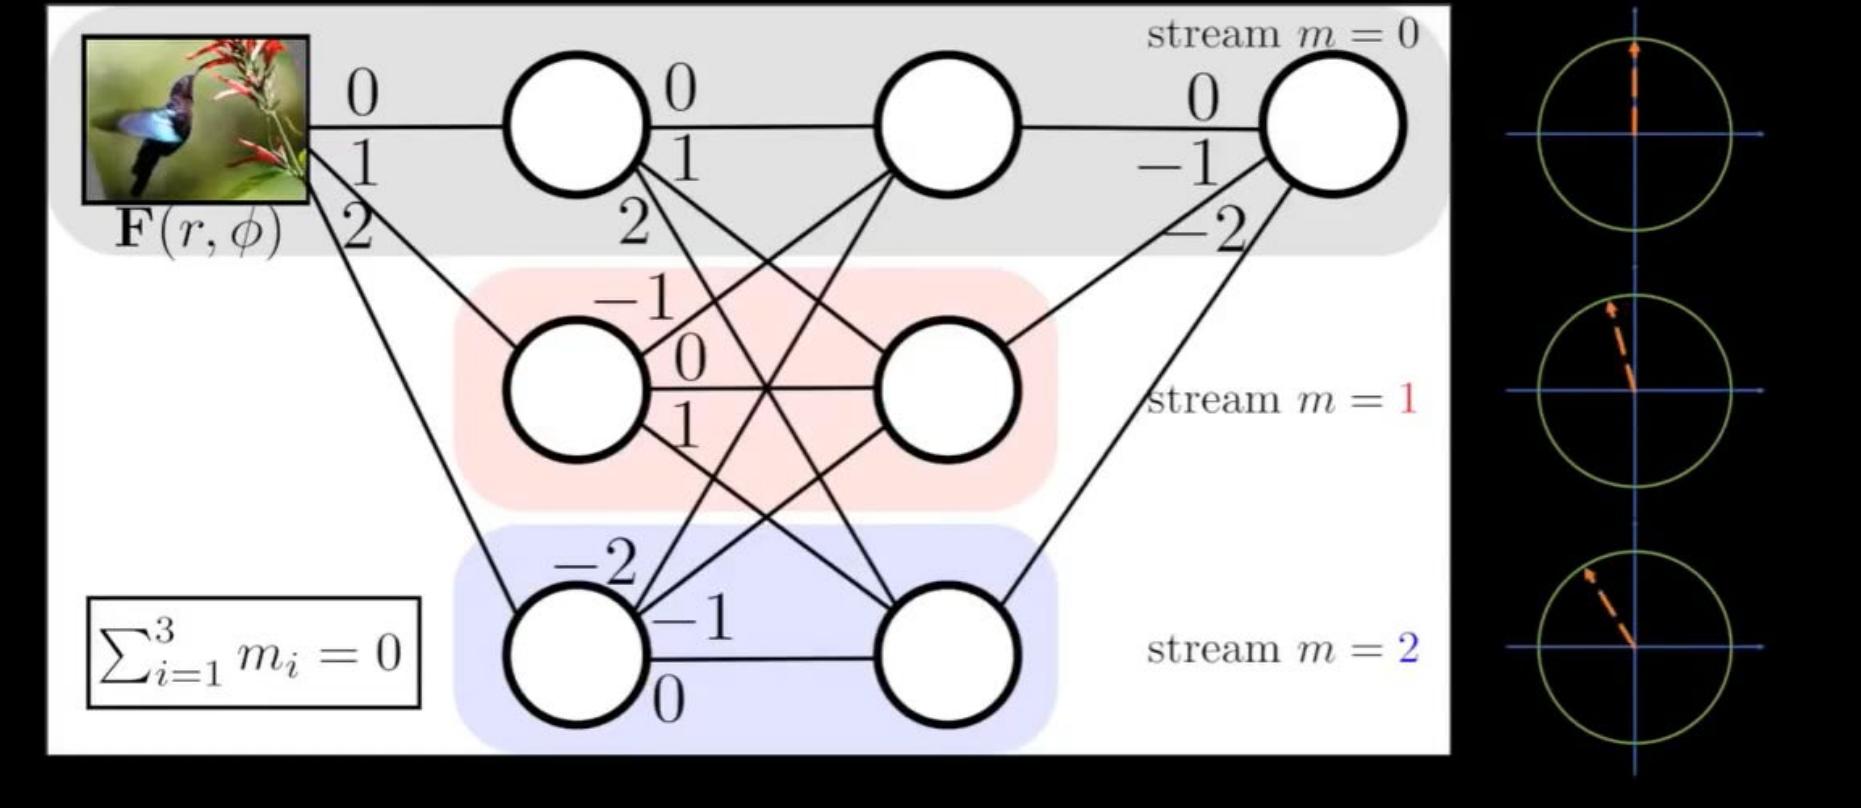
*



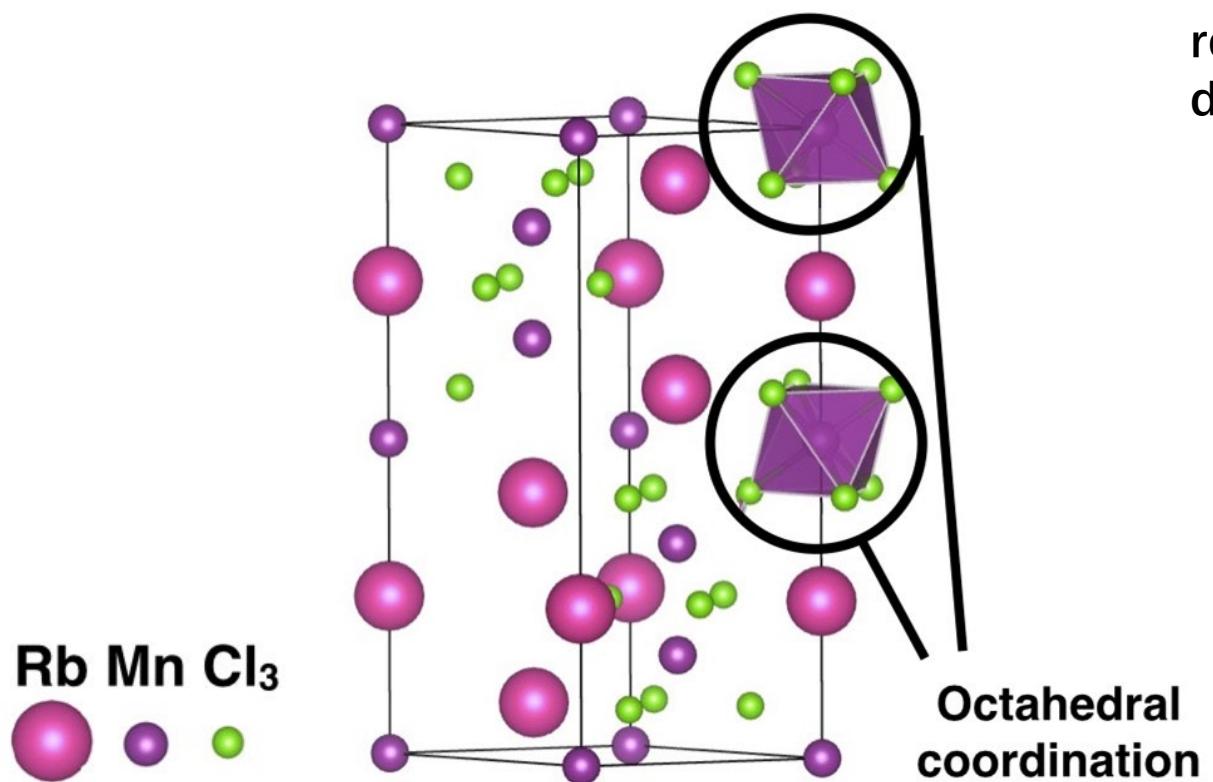
Length = intensity

CONVOLUTIONS ON ROTO-TRANSLATION-Harmonic Networks

Harmonic Networks

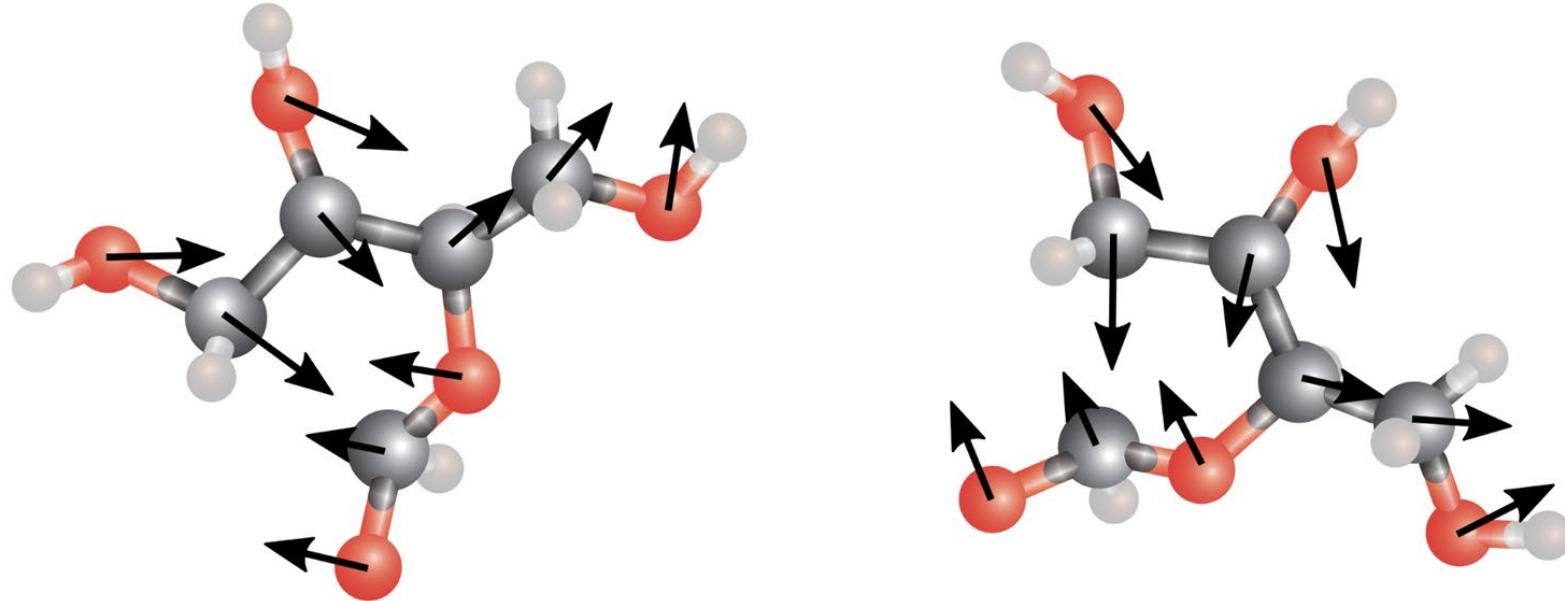


What assumptions do we want “built in” to our neural networks (for materials data)?



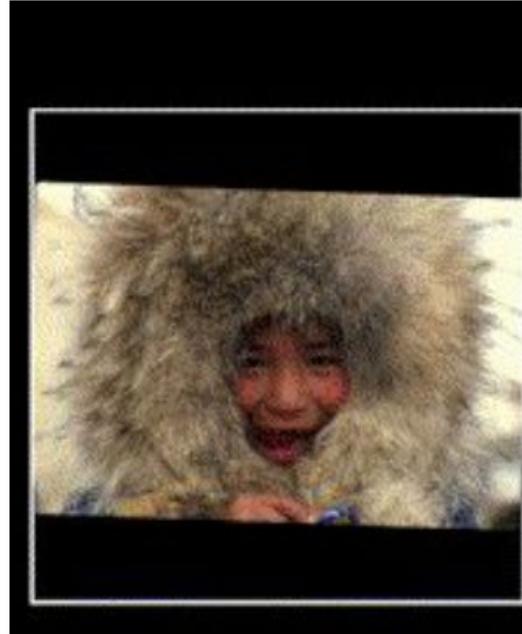
These networks can recognize **equivalent** recurring **geometric** patterns that appear in different **locations** and **orientations**.

Given a molecule and a rotated copy,
predicted forces are the same up to rotation.
(Predicted forces are equivariant to rotation.)



Previous work on 2D
rotation-equivariance uses
filters based on **circular**
harmonics.

To extend to 3D, we use
spherical harmonics.



Rotated image



CNN filter output



Harmonic filter output

<http://visual.cs.ucl.ac.uk/pubs/harmonicNets/>

Apr 2017

Harmonic Networks: Deep Translation and Rotation Equivariance

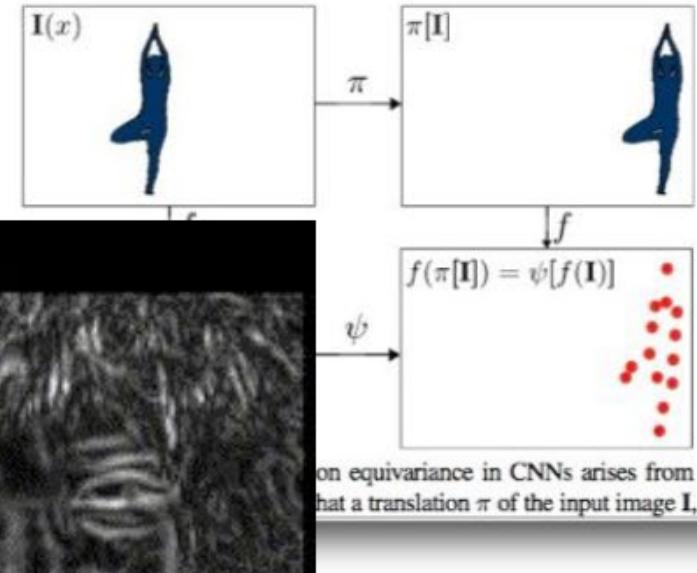
Daniel E. Worrall, Stephan J. Garbin, Daniyar Turmukhambetov and Gabriel J. Brostow

{d.worrall, s.garbin, d.turmukhambetov, g.brostow}@cs.ucl.ac.uk

University College London*

Abstract

Translating or rotating an input image should not affect the results of many computer vision tasks. Convolutional neural networks (CNNs) are already translation equivariant: input image translations produce proportionate feature map translations.



The input and output of our network is represented as tensors with point (or atom), channel, and representation indices organized by irreducible representation.

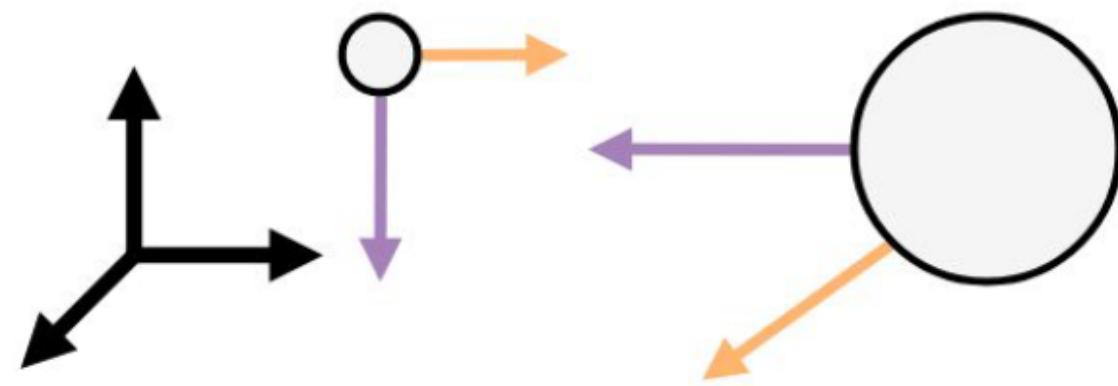
$$V(l)_{acm} = \{ 0: [[m0]], [[m1]] , \\ 1: [[v0x, v0y, v0z], [a0x, a0y, a0z]], \\ [[v1x, v1y, v1z], [a1x, a1y, a1z]] \}$$

l : dictionary key, l

$[]$: point index, a

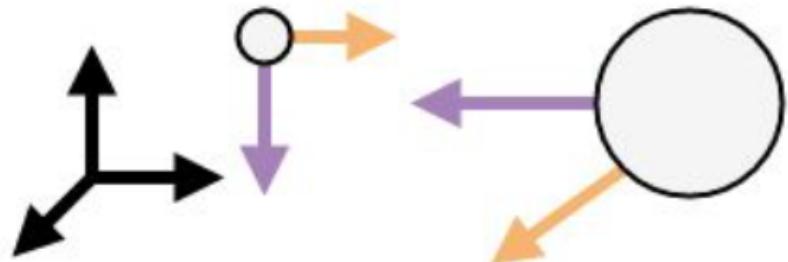
$[]$: channel index, c

$[]$: representation index, m



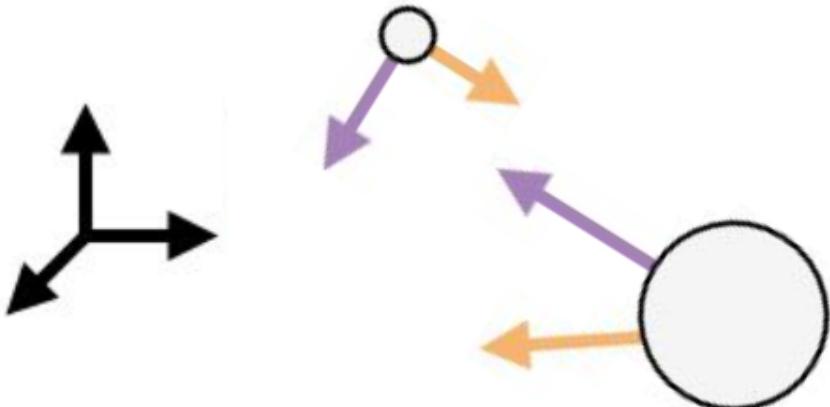
Geometric tensors transform predictably under 3D rotation.

Two point **masses** with **velocity** and **acceleration**.



```
geometry = [[x0, y0, z0], [x1, y1, z1]]  
features = [  
    [m0, v0y, v0z, v0x, a0y, a0z, a0x]  
    [m1, v1y, v1z, v1x, a1y, a1z, a1x]]
```

Same system, with rotated coordinates.



Convolutional kernels...

Learned
Parameters

with no symmetry:

$$W(\vec{r})$$

with 2D circular harmonics:

$$R(r) e^{im\phi} \longrightarrow [-\pi, \pi]$$

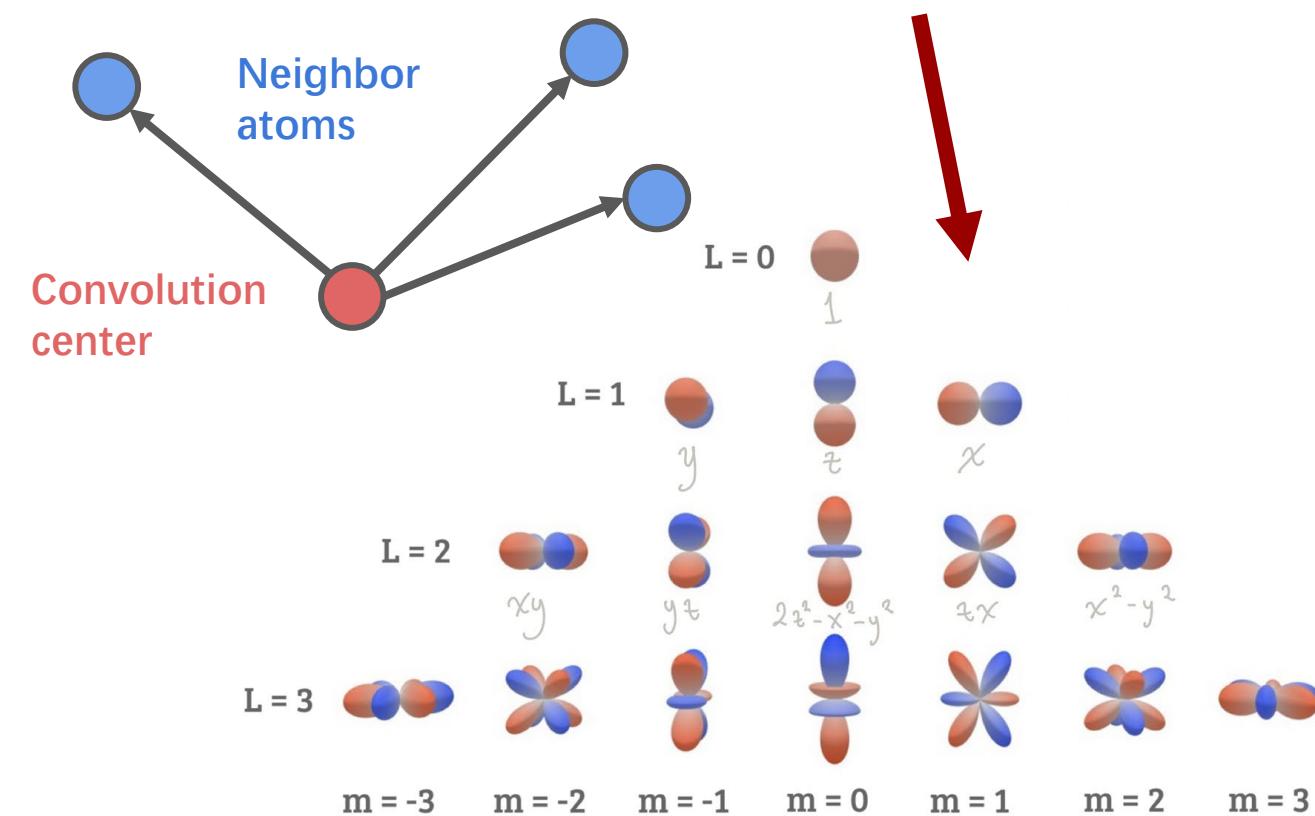
with 3D spherical harmonics:

$$R(r) Y_l^m(\hat{r})$$

Euclidean Neural Networks are similar to convolutional neural networks...

Equivariant convolutional filters are based
on learned radial functions and spherical
harmonics...

$$W(\vec{r}) = R(r)Y_l^m(\hat{r})$$



	$L = 0$	
	$L = 1$	
	$L = 2$	
		Irreducible representations
Scalars fields	$l = 0$	
Vectors fields	$l = 1$	
3x3 Matrix fields	$l = 0 \oplus 1 \oplus 2$	

```

Rs_s_orbital = o3.Irrep("0e")

Rs_p_orbital = o3.Irrep("1o")

Rs_d_orbital = o3.Irrep("2e")

Rs_f_orbital = o3.Irrep("3o")

scalar = e3nn.o3.Irrep("0e") # L=0, even
vector = e3nn.o3.Irrep("1o") # L=1, odd

irreps = 1 * scalar + 1 * vector + 1 *
vector

```

Interactions in **equivariant models** are more complex.

How do we interact **invariant** objects? Scalar multiplication.

$$\begin{array}{c} \textcolor{purple}{\blacksquare} \\ \times \\ \textcolor{orange}{\blacksquare} \end{array} = \begin{array}{c} \textcolor{red}{\blacksquare} \end{array}$$

How do we interact **equivariant** objects?
Tensor products!

$$\begin{array}{c} \nearrow \\ \textcolor{red}{\nearrow} \\ \otimes \\ \textcolor{red}{\searrow} \\ \searrow \end{array} = \begin{array}{c} \textcolor{red}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \quad \textcolor{red}{\blacksquare} \\ \textcolor{orange}{\blacksquare} \quad \textcolor{orange}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \\ \textcolor{purple}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \quad \textcolor{red}{\blacksquare} \end{array} = \begin{array}{c} \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{red}{\blacksquare} \end{array} \otimes \begin{array}{c} \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{orange}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{orange}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \end{array} \otimes \begin{array}{c} \textcolor{purple}{\blacksquare} \quad \textcolor{red}{\blacksquare} \quad \textcolor{red}{\blacksquare} \\ \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \end{array}$$

Generalizes to higher orders. Same mathematics that describes atomic interactions, e.g. selection rules in spectroscopy.

$$\begin{array}{c} \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{red}{\blacksquare} \end{array} \otimes \begin{array}{c} \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{orange}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \\ \textcolor{orange}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \end{array} \otimes \begin{array}{c} \textcolor{purple}{\blacksquare} \quad \textcolor{red}{\blacksquare} \quad \textcolor{red}{\blacksquare} \\ \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{white}{\blacksquare} \\ \textcolor{red}{\blacksquare} \quad \textcolor{white}{\blacksquare} \quad \textcolor{purple}{\blacksquare} \end{array}$$

dot product
trace
invariant
 $L=0$
1 degree of freedom

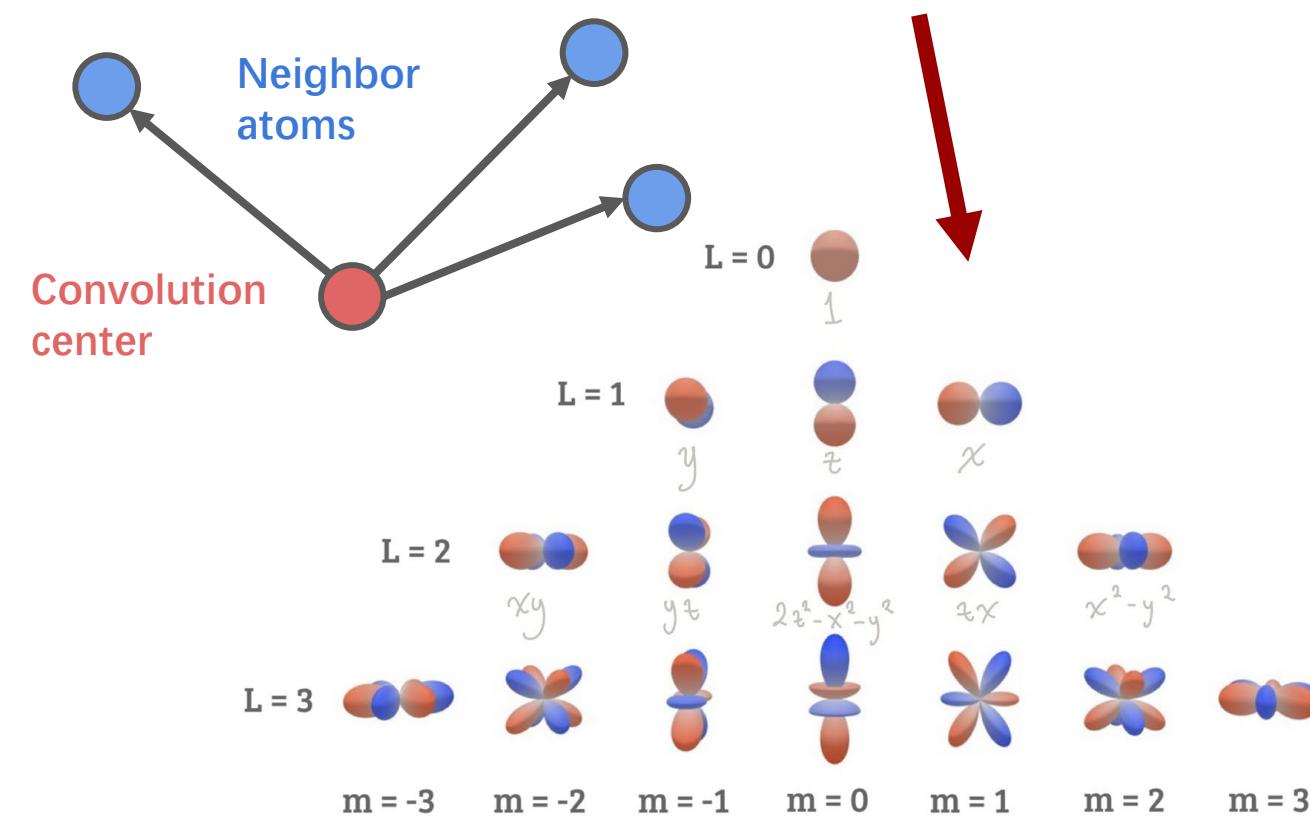
cross-product
antisymmetric
equivariant
 $L=1$
3 degrees of freedom

symmetric
traceless
equivariant
 $L=2$
5 degrees of freedom

Euclidean Neural Networks are similar to convolutional neural networks...

Equivariant convolutional filters are based
on learned radial functions and spherical
harmonics...

$$W(\vec{r}) = R(r)Y_l^m(\hat{r})$$



Spherical harmonics

Y_l^m

$L = 0$



angular portion of
hydrogenic wavefunctions

$L = 1$



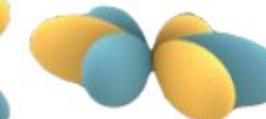
basis functions for $(2l + 1)$
dimensional irreducible
representations of $SO(3)$

basis functions for signals
on a sphere

$L = 2$



$L = 3$



$m = -3$

$m = -2$

$m = -1$

$m = 0$

$m = 1$

$m = 2$ $m = 3$

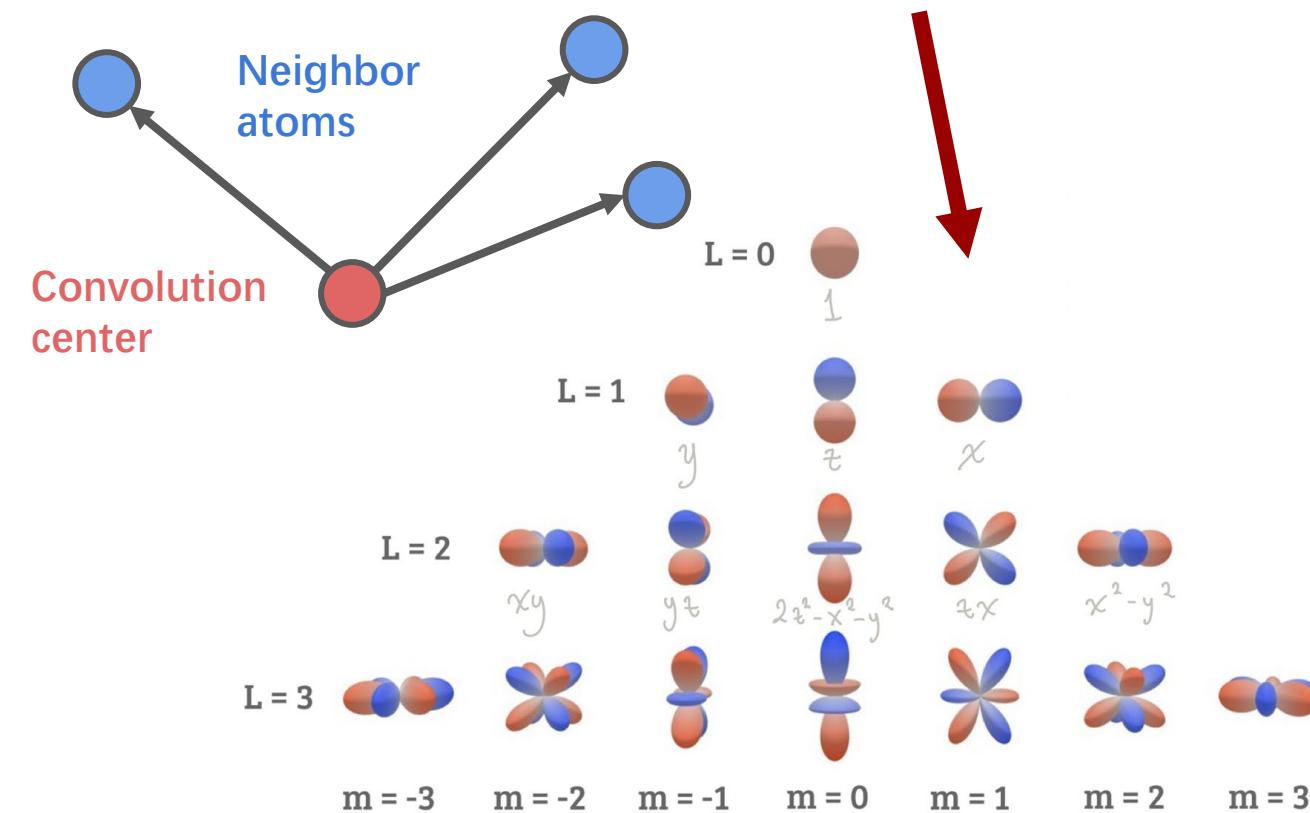
Inigo.quilez

https://en.wikipedia.org/wiki/Spherical_harmonics

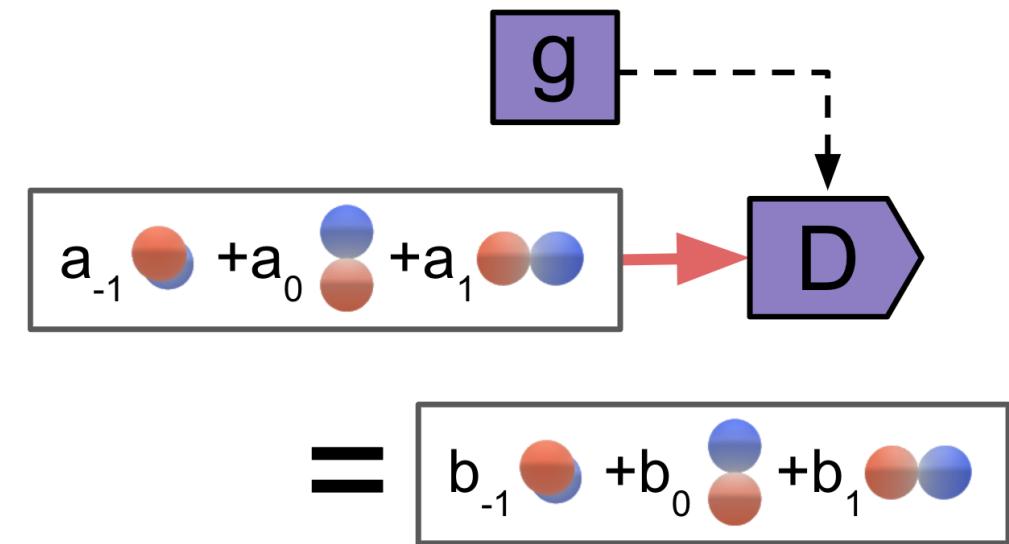
Euclidean Neural Networks are similar to convolutional neural networks...

Equivariant convolutional filters are based on learned radial functions and spherical harmonics...

$$W(\vec{r}) = R(r)Y_l^m(\hat{r})$$



Spherical harmonics of the same L transform together under rotation g .



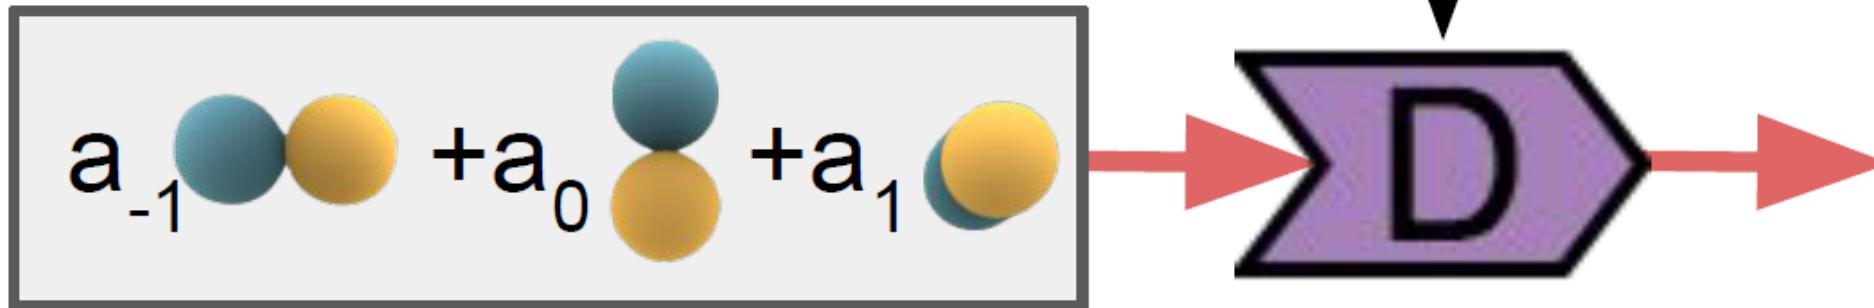
Spherical harmonics transform in the same manner as the irreducible representations of $SO(3)$.

Spherical harmonics of a given L transform together under rotation.

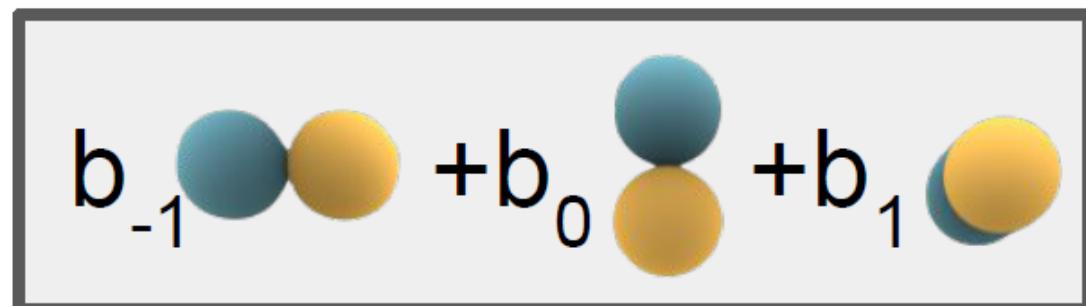
Let g be a 3d rotation matrix.



D is the Wigner-D matrix.
It has shape $[2l + 1, 2l + 1]$
and is a function of g .



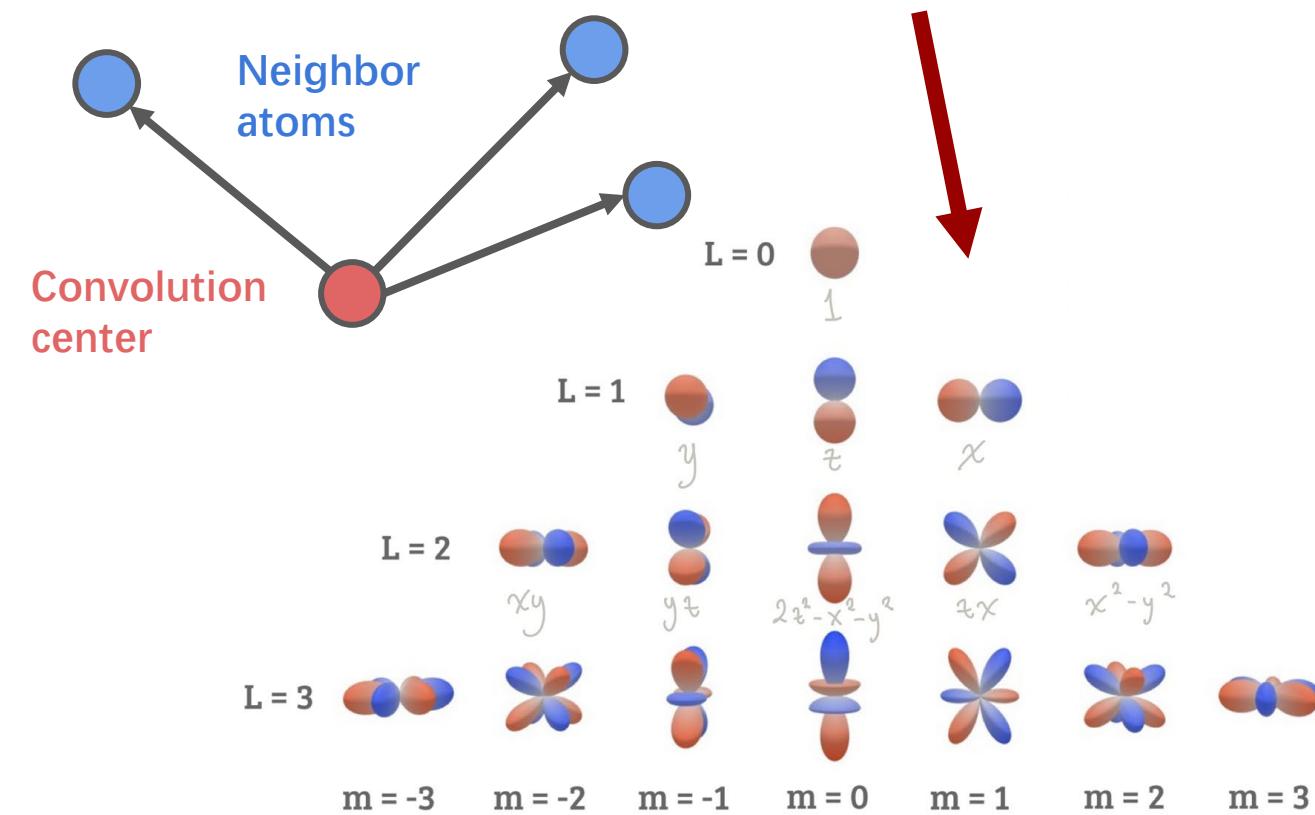
=



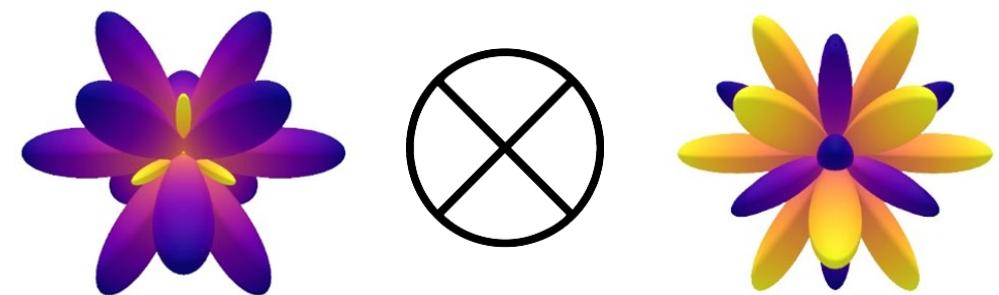
Euclidean Neural Networks are similar to convolutional neural networks...

Equivariant convolutional filters are based on **learned radial functions** and **spherical harmonics**...

$$W(\vec{r}) = R(r)Y_l^m(\hat{r})$$



...and **geometric tensor algebra** allow us to generalize scalar operations to more complex geometric tensors.



e.g. How to multiply two vectors?

$$\vec{a} \cdot \vec{b} = c \quad \text{scalar}$$

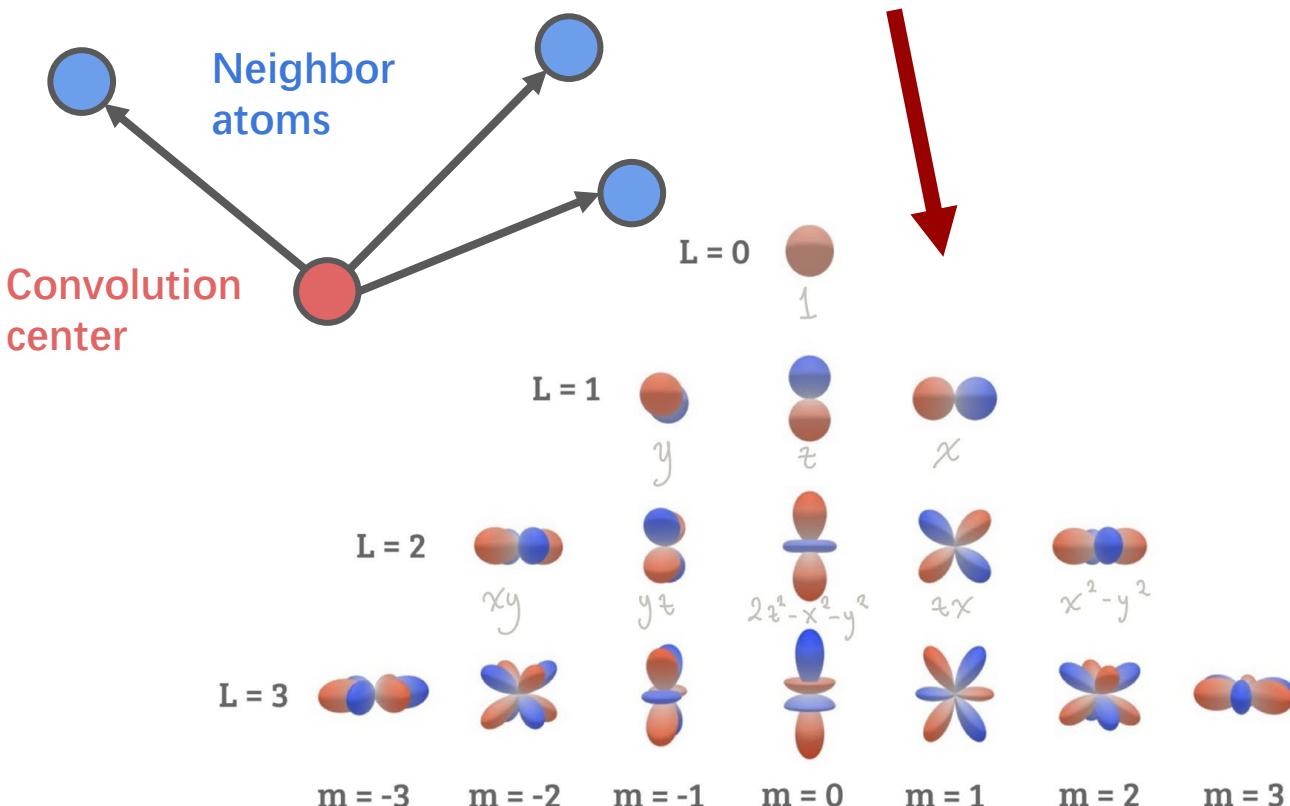
$$\vec{a} \times \vec{b} = \vec{c} \quad \text{vector}$$

$$\vec{a} \otimes \vec{b} = C \quad 3 \times 3 \text{ matrix}$$

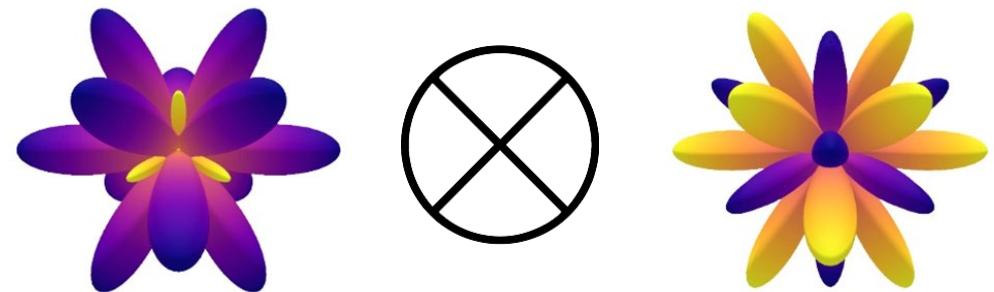
Euclidean Neural Networks are similar to convolutional neural networks...

Equivariant convolutional filters are based on learned radial functions and spherical harmonics...

$$W(\vec{r}) = R(r) Y_l^m(\hat{r})$$



...and **geometric tensor algebra** allow us to generalize scalar operations to more complex geometric tensors.



34. CLEBSCH-GORDAN COEFFICIENTS, SPHERICAL HARMONICS, AND d FUNCTIONS

Note: A square-root sign is to be understood over every coefficient, e.g., for $-8/15$ read $-\sqrt{8/15}$

Notation:

J	J	\dots
M	M	\dots
m_1	m_2	\dots
m_1	m_2	Coefficients
\vdots	\vdots	\vdots

Y₁⁰ = $\sqrt{\frac{3}{4\pi}} \cos \theta$

Y₁¹ = $-\sqrt{\frac{3}{8\pi}} \sin \theta e^{i\phi}$

Y₂⁰ = $\sqrt{\frac{5}{16\pi}} \left(\frac{3}{2} \cos^2 \theta - \frac{1}{2} \right)$

Y₂¹ = $-\sqrt{\frac{15}{32\pi}} \sin \theta \cos \theta e^{i\phi}$

Y₂² = $\frac{1}{4} \sqrt{\frac{15}{16\pi}} \sin^2 \theta e^{2i\phi}$

Y₃⁰ = $\sqrt{\frac{5}{64\pi}} \left(\frac{5}{2} \cos^3 \theta - \frac{3}{2} \cos \theta + \frac{1}{2} \right)$

Y₃¹ = $-\sqrt{\frac{15}{128\pi}} \sin \theta \left(\frac{5}{2} \cos^2 \theta - \frac{1}{2} \right) e^{i\phi}$

Y₃² = $\sqrt{\frac{15}{128\pi}} \sin^2 \theta \cos \theta e^{2i\phi}$

Y₃³ = $-\sqrt{\frac{15}{256\pi}} \sin^3 \theta e^{3i\phi}$

$Y_\ell^{-m} = (-1)^m Y_\ell^m$

$d_{m,0}^\ell = \sqrt{\frac{4\pi}{2\ell+1}} Y_\ell^m e^{-im\phi}$

$(j_1 j_2 m_1 m_2 | j_1 j_2 JM) = (-1)^{-j_1-j_2} (j_2 j_1 m_2 m_1 | j_2 j_1 JM)$

The input and output of our network is represented as tensors with point (or atom), **channel**, and **representation** indices organized by irreducible representation.

$$V(l)_{acm} = \text{IN} \quad \begin{array}{c} \text{IN} \\ \swarrow \quad \searrow \\ \text{Red} \quad \text{Grey} \quad \text{Green} \end{array}$$

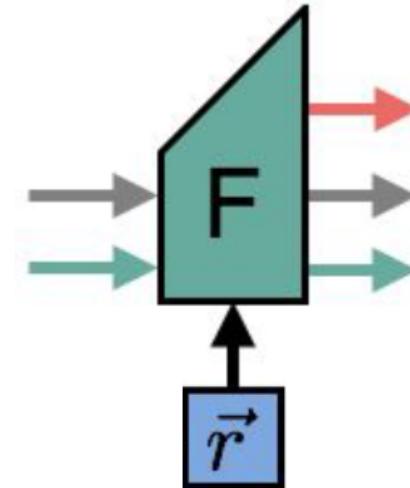
Representation →

Points →

Channels →

Filters contribute a **representation** index due to use of spherical harmonics.

$$R(r)Y_l^m(\hat{r})$$



Representation →

Points →

Channels →

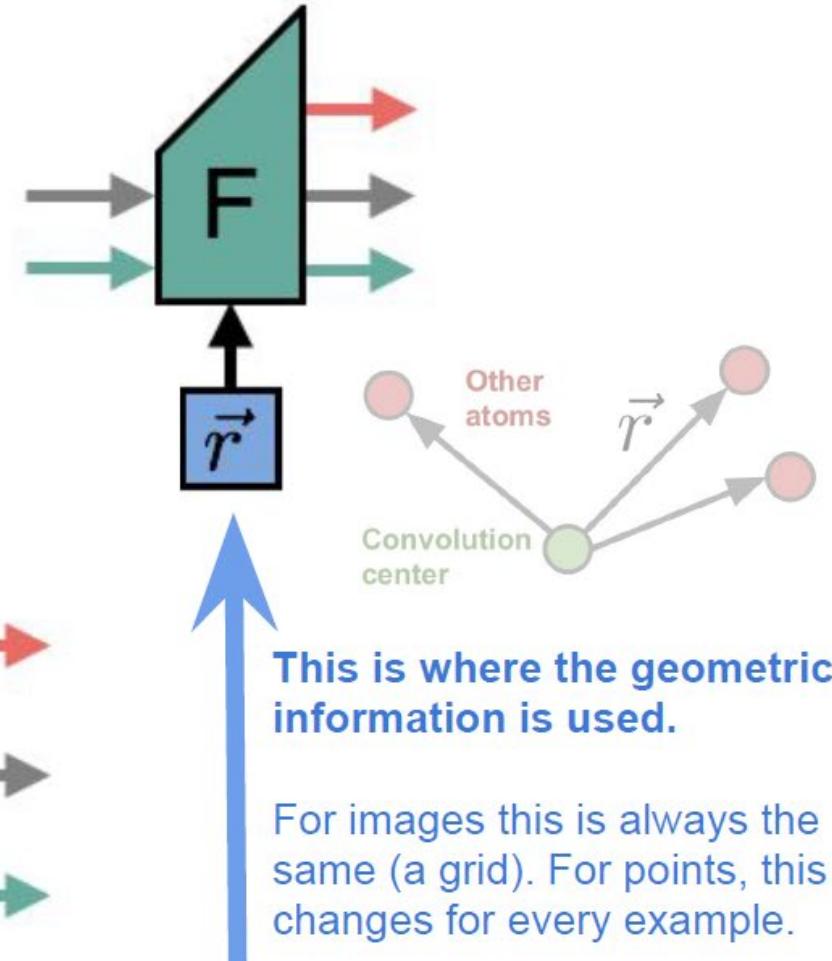
Filters contribute a **representation** index due to use of spherical harmonics.

$$R(r)Y_l^m(\hat{r})$$

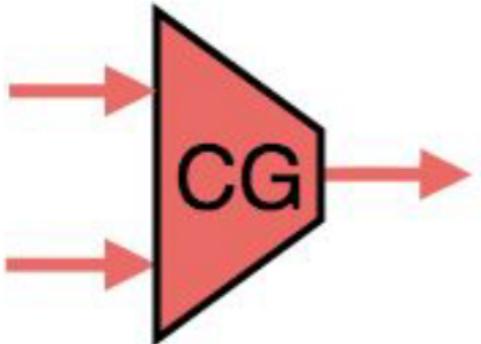
Representation

Points

Channels



To combine two tensors to create one tensor, we use Clebsch-Gordan coefficients.



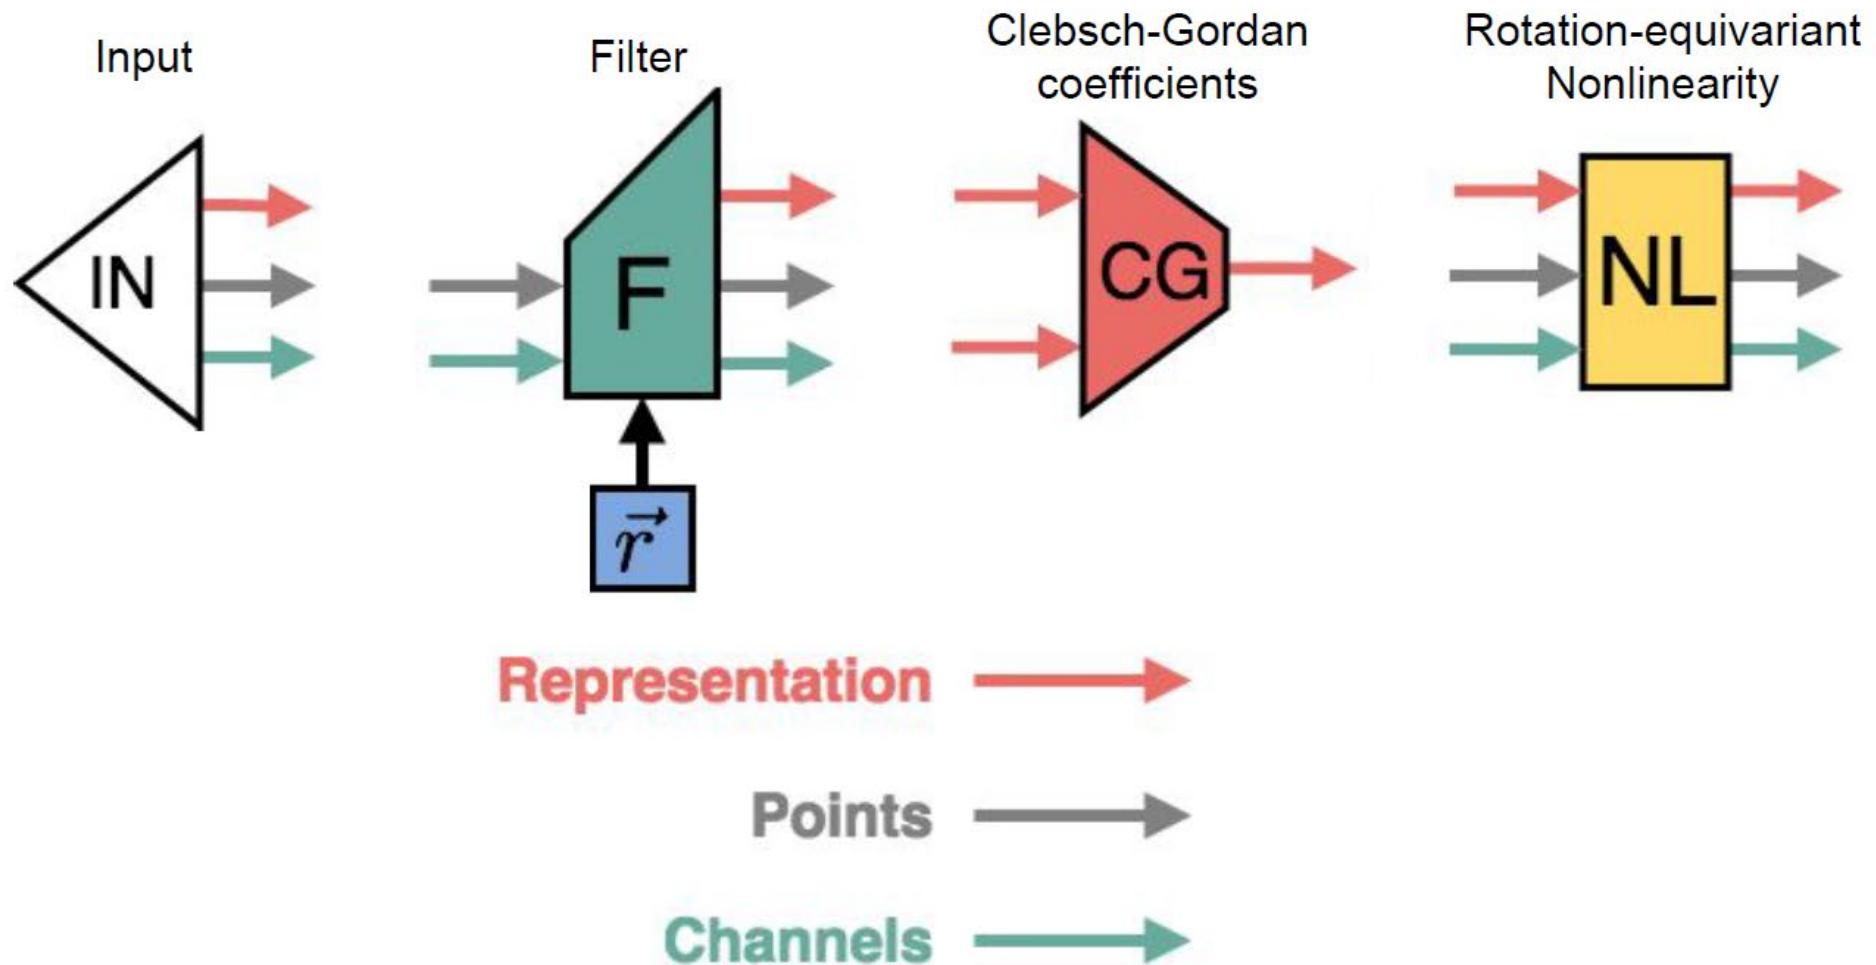
Representation →
Points →
Channels →

34. CLEBSCH-GORDAN COEFFICIENTS, SPHERICAL HARMONICS, AND d FUNCTIONS

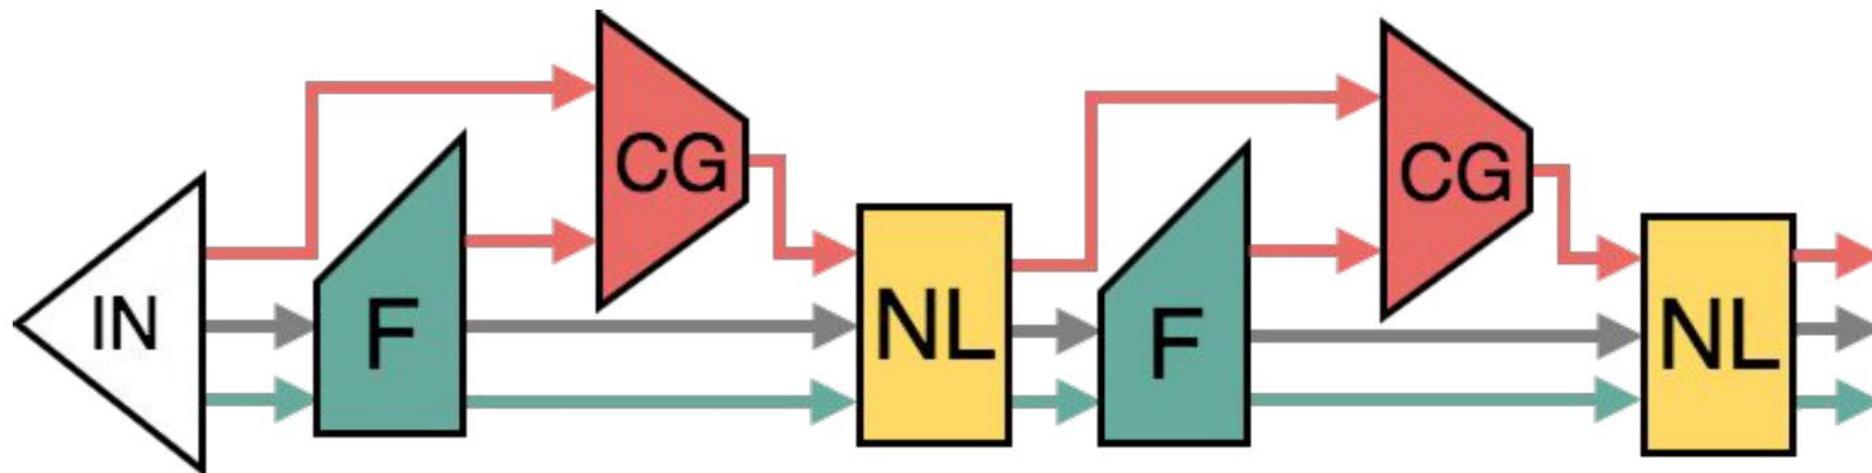
Note: A square-root sign is to be understood over every coefficient, e.g., for $-8/15$ read $-\sqrt{8/15}$

J_1	J_2	M_1	M_2	Notation:
m_1	m_2	m_1	m_2	Coefficients
.
$Y_1^0 = \sqrt{\frac{3}{4\pi}} \cos \theta$	$2 \times 1/2$	$\begin{matrix} 5/2 \\ +5/2 \\ +2 \\ +1/2 \end{matrix}$	$\begin{matrix} 5/2 \\ 5/2 \\ 1/5 \\ 4/5 \end{matrix}$	$\begin{matrix} J \\ M \\ M \\ M \end{matrix}$
$Y_1^1 = -\sqrt{\frac{3}{8\pi}} \sin \theta e^{i\phi}$	$+2$	$\begin{matrix} 1/2 \\ +1/2 \\ +1/2 \\ +1/2 \end{matrix}$	$\begin{matrix} 1/5 \\ 4/5 \\ 5/2 \\ 3/2 \end{matrix}$	
$Y_2^0 = \sqrt{\frac{5}{4\pi}} \left(\frac{3}{2} \cos^2 \theta - \frac{1}{2} \right)$	$+3/2$	$\begin{matrix} 1/2 \\ +1/2 \\ +1/2 \\ +1/2 \end{matrix}$	$\begin{matrix} 5/2 \\ 3/2 \\ 3/5 \\ -2/5 \end{matrix}$	
$Y_2^1 = -\sqrt{\frac{15}{8\pi}} \sin \theta \cos \theta e^{i\phi}$	$+1/2$	$\begin{matrix} 0 \\ 1/2 \\ 1/2 \\ 1/2 \end{matrix}$	$\begin{matrix} 0 \\ -1/2 \\ 3/5 \\ 2/5 \end{matrix}$	$\begin{matrix} J \\ M \\ M \\ M \end{matrix}$
$Y_2^2 = \frac{1}{4} \sqrt{\frac{15}{2\pi}} \sin^2 \theta e^{2i\phi}$	$3/2 \times 1/2$	$\begin{matrix} 2 \\ +2 \\ +1/2 \end{matrix}$	$\begin{matrix} 2 \\ 2 \\ 1 \end{matrix}$	
$3/2 \times 1/2$	$+1/2$	$\begin{matrix} 1/2 \\ +1/2 \\ +1/2 \end{matrix}$	$\begin{matrix} 1/2 \\ 1/2 \\ 1/2 \end{matrix}$	$\begin{matrix} J \\ M \\ M \\ M \end{matrix}$
$1 \times 1/2$	$+3/2$	$\begin{matrix} 3/2 \\ +3/2 \\ +1/2 \end{matrix}$	$\begin{matrix} 3/2 \\ 1/2 \\ 1/2 \end{matrix}$	
$1/2 \times 1/2$	$+1/2$	$\begin{matrix} 1/2 \\ +1/2 \\ +1/2 \end{matrix}$	$\begin{matrix} 1/2 \\ 1/2 \\ 1/2 \end{matrix}$	
1×1	$+3$	$\begin{matrix} 3 \\ +2 \\ +1 \end{matrix}$	$\begin{matrix} 1 \\ -1/2 \\ 1 \end{matrix}$	
1×1	$+2$	$\begin{matrix} 2 \\ +2 \\ +1 \end{matrix}$	$\begin{matrix} 2 \\ -1/2 \\ 1 \end{matrix}$	
1×1	$+1$	$\begin{matrix} 1 \\ +1 \\ +1 \end{matrix}$	$\begin{matrix} 1 \\ 0 \\ 0 \end{matrix}$	
$Y_\ell^{-m} = (-1)^m Y_\ell^{m*}$	0	$\begin{matrix} 0 \\ 0 \\ 0 \end{matrix}$	$\begin{matrix} 0 \\ 0 \\ 0 \end{matrix}$	
	$d_{m,0}^\ell = \sqrt{\frac{4\pi}{2\ell+1}} Y_\ell^m e^{-im\phi}$	$\begin{matrix} 0 \\ -1 \\ -1 \end{matrix}$	$\begin{matrix} 0 \\ 1 \\ 1 \end{matrix}$	$(j_1 j_2 m_1 m_2 j_1 j_2 JM)$
		$\begin{matrix} 1/2 \\ 1/2 \\ 1/2 \end{matrix}$	$\begin{matrix} 1/2 \\ 1/2 \\ 1/2 \end{matrix}$	$= (-1)^{J-j_1-j_2} (j_2 j_1 m_2 m_1 j_2 j_1 JM)$

These are components of **tensor field networks**



This is what a two-layer tensor field network looks like:



Representation →

Points →

Channels →

... and this is what one layer (without NL) looks like with all the indices written out:

$$\begin{aligned} \mathcal{L}_{acm_O}^{(l_O)}(\vec{r}_a, V_{acm_I}^{(l_I)}) \\ := \sum_{m_F, m_I} C_{(l_F, m_F)(l_I, m_I)}^{(l_O, m_O)} \sum_{b \in S} F_{cm_F}^{(l_F, l_I)}(\vec{r}_{ab}) V_{bcm_I}^{(l_I)} \end{aligned}$$

... and this is what one layer (without NL) looks like with all the indices written out:

$$\begin{aligned}\mathcal{L}_{acmo}^{(l_O)}(\vec{r}_a, V_{acm_I}^{(l_I)}) \\ := \sum_{m_F, m_I} C_{(l_F, m_F)(l_I, m_I)}^{(l_O, m_O)} \sum_{b \in S} F_{cm_F}^{(l_F, l_I)}(\vec{r}_{ab}) V_{bcm_I}^{(l_I)}\end{aligned}$$

$$F_{cm_F}^{(l_F, l_I)}(\vec{r}_{ab}) = R_c^{(l_F, l_I)}(r_{ab}) Y_{m_F}^{(l_F)}(\hat{r}_{ab})$$

... and this is what one layer (without NL) looks like with all the indices written out:

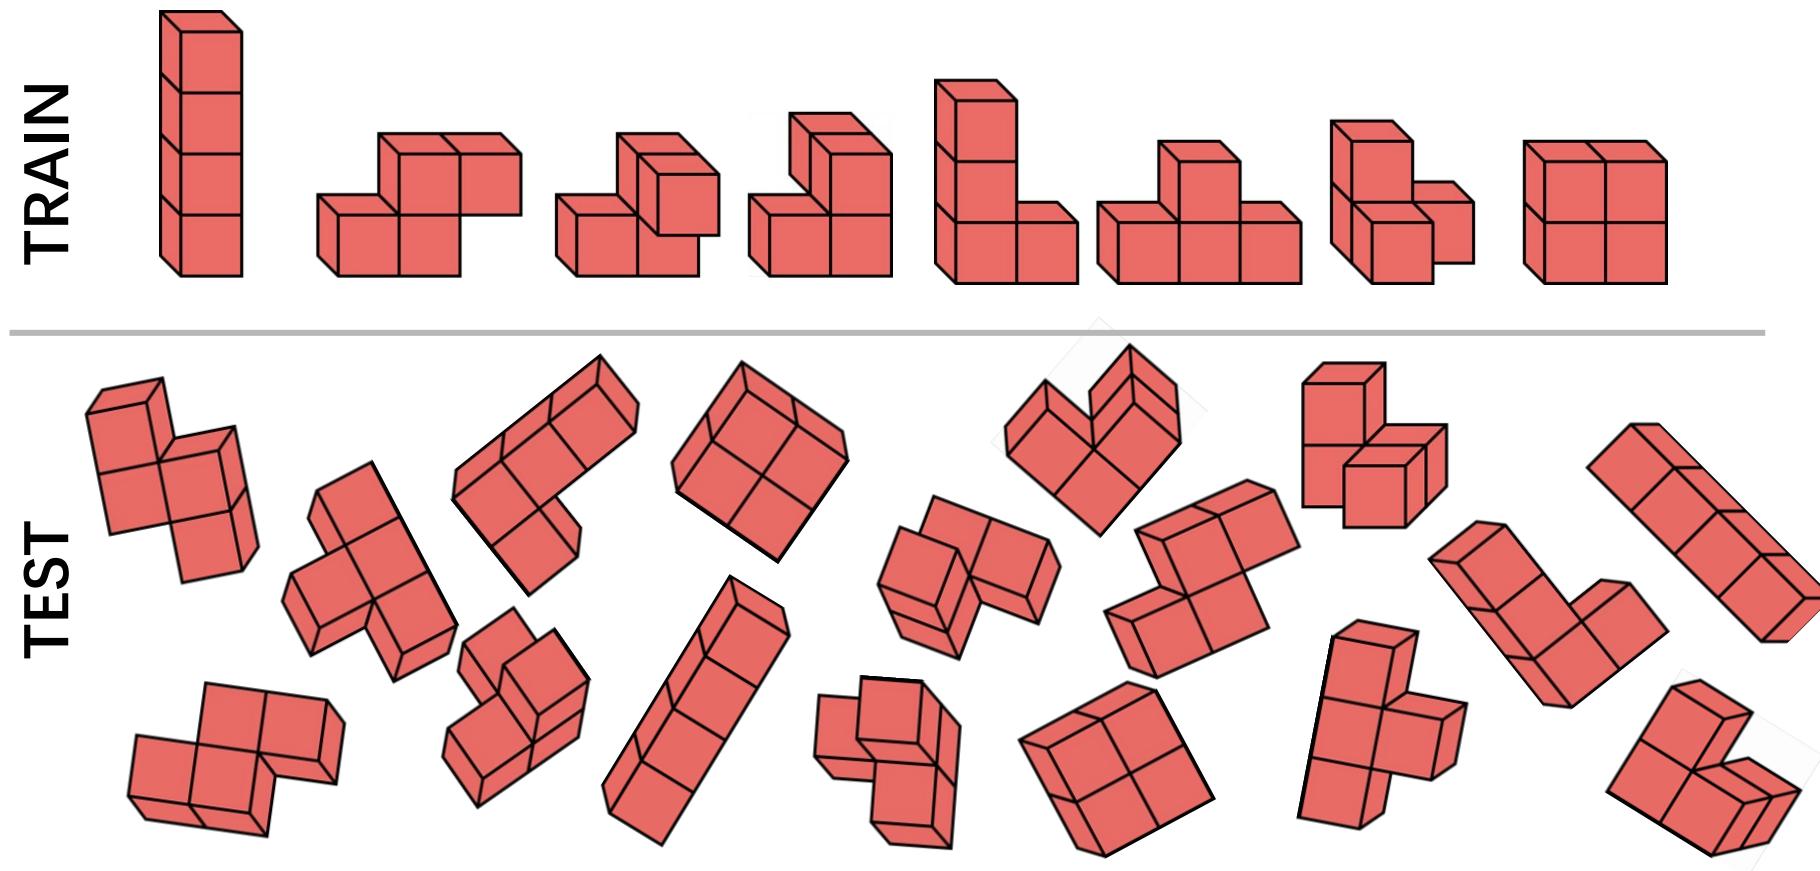
$$\mathcal{L}_{acm_O}^{(l_O)}(\vec{r}_a, V_{acm_I}^{(l_I)})$$

$$:= \sum_{m_F, m_I} C_{(l_F, m_F)(l_I, m_I)}^{(l_O, m_O)} \sum_{b \in S} F_{cm_F}^{(l_F, l_I)}(\vec{r}_{ab}) V_{bcm_I}^{(l_I)}$$

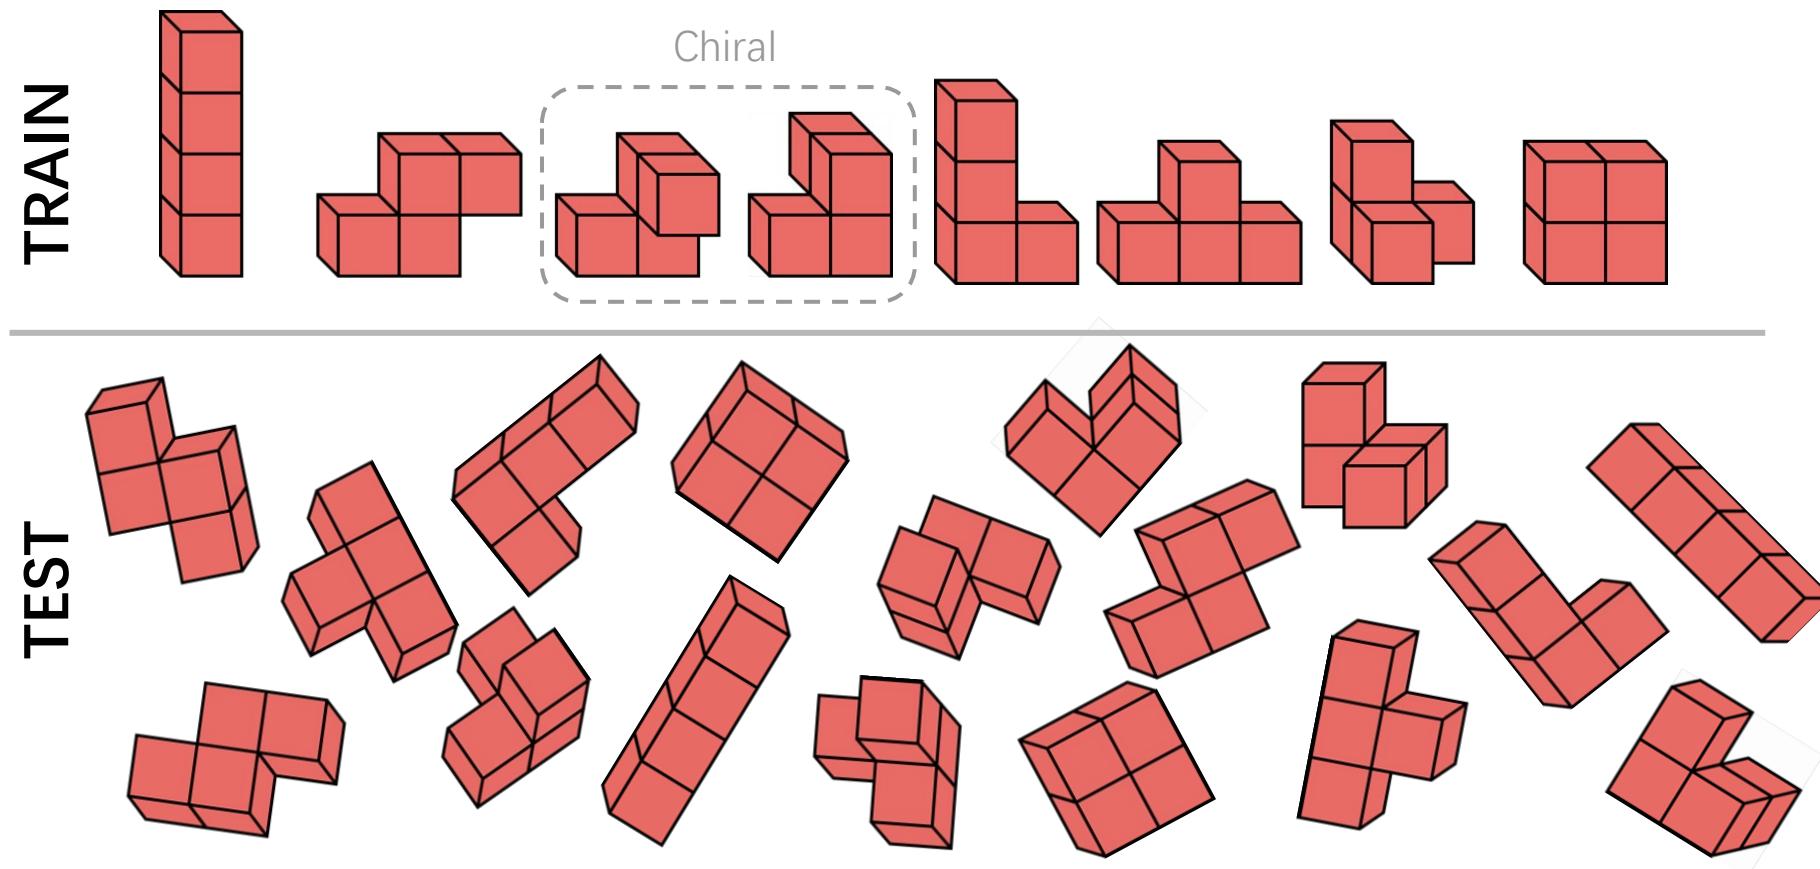
$$F_{cm_F}^{(l_F, l_I)}(\vec{r}_{ab}) = R_c^{(l_F, l_I)}(r_{ab}) Y_{m_F}^{(l_F)}(\hat{r}_{ab})$$

$$R_c^{(l_F, l_I)}(r_{ab}) = \sum_h W_{ch} \text{NL}\left(\sum_j W_{hj} B_j(r_{ab}) + b_h\right) + b_c$$

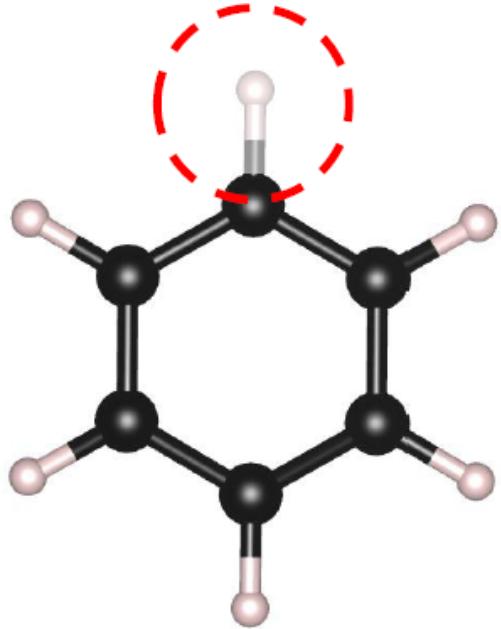
Our unit test: Trained on 3D Tetris shapes in one orientation,
these network can perfectly identify these shapes in any orientation.



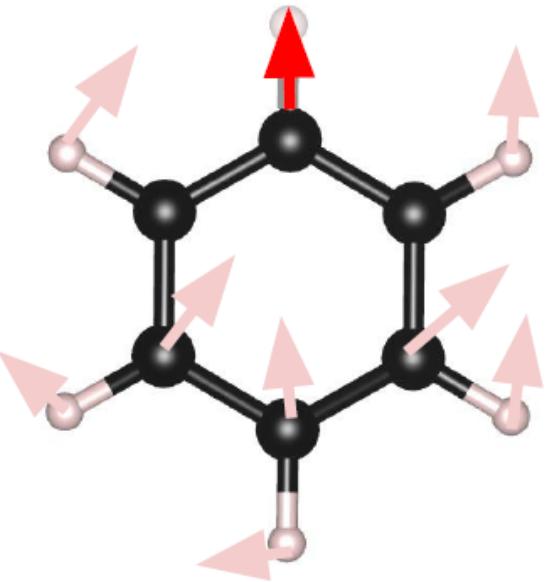
Our unit test: Trained on 3D Tetris shapes in one orientation,
these network can perfectly identify these shapes in any orientation.



Given a small organic molecule with an atom removed, replace the correct element at the correct location in space.



Input coordinates with missing atom.



Network outputs
(N-1) atom type features (scalars),
(N-1) displacement vectors, and
(N-1) scalars indicating confidence probability used for "voting".

DATASET

QM9: <http://www.quantum-machine.org/datasets/>
134k molecules with 9 or less heavy atoms (non-hydrogen) and elements H, C, N, O, F.

TRAIN

1,000 molecules with 5-18 atoms

TEST

1,000 molecules with 19 atoms
1,000 molecules with 23 atoms
1,000 molecules with 25-29 atoms

Atoms	Number of predictions	Accuracy (%) ($\leq 0.5 \text{ \AA}$ and atom type)	Distance MAE in \AA
5-18 (train)	15 947	92.6	0.16
19	19 000	94.7	0.15
23	23 000	96.9	0.14
25-29	25 404	97.8	0.17

Learns to replace atoms with over 90% accuracy across train and test by seeing the same 1,000 molecules 200 times.