

Latent Collaboration in Multi-Agent Systems

Jiaru Zou Xiyuan Yang Ruizhong Qiu Gaotang Li
Katherine Tieu Pan Lu Ke Shen Hanghang Tong
Yejin Choi Jingrui He James Zou Mengdi Wang Ling
Yang

Presenter: Yicheng Tao

December 15, 2025

Outline

Introduction

Method

Experiment

Discussion

Introduction

Method

Experiment

Discussion

Overview

TLDR

A comprehensive model collaboration framework unifying both latent reasoning and latent communication.

Summary

- ▶ **Latent Reasoning:** Leverage hidden representations within transformers to enable single model's internal latent chain-of-thought reasoning.
- ▶ **Latent Communication:** Employ KV caches for information exchange across two models.

Key Feature

- ▶ Training-free
- ▶ Lower computational complexity
- ▶ Higher performance (Need check)

Preliminary

Notation

Let W_{in} denote the input embedding layer, W_{out} denote the language model head, $E = [e_1, e_2, \dots, e_t] \in \mathbb{R}^{t \times d_h}$ denote the input token embeddings and $H = [h_1, h_2, \dots, h_t] \in \mathbb{R}^{t \times d_h}$ denote the final-layer hidden representations.

KV Cache

In decoder-only Transformers, the Key-Value (KV) cache functions as a **dynamic working memory** during auto-regressive generation, storing intermediate representations from previous decoding steps to avoid redundant computation. Specifically, when the next token at step $t + 1$ is generated, the KV cache at each layer ($K_{\text{cache}}, V_{\text{cache}}$) will be updated as:

$$K_{\text{cache}} \leftarrow [K_{\leq t}; K_{t+1}], \quad V_{\text{cache}} \leftarrow [V_{\leq t}; V_{t+1}]$$

Introduction

Method

Experiment

Discussion

System Overview

Latent Collaboration (LatentMAS)

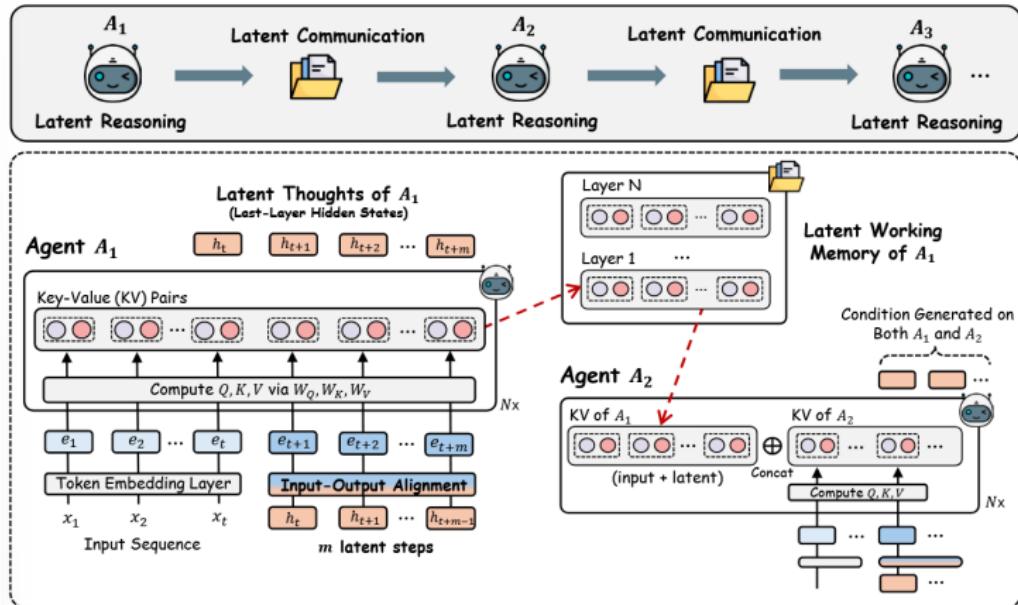


Figure 3 | Overview of LatentMAS. Each LLM agent in the system first generates latent thoughts through last-layer hidden states, then transfers information layer-wise via shared latent working memory stored in KV-caches, enabling completely system-wide latent collaboration.

Latent Reasoning

Generation Procedure

Given the input embeddings $E = [e_1, e_2, \dots, e_t]$ and the obtained last-layer hidden representation h_t at step t , insert h_t into the input embedding instead of decoding and next-token embedding. m latent steps yield a sequence of newly generated last-layer hidden states $H = [h_{t+1}, h_{t+2}, \dots, h_{t+m}]$, which is defined as **latent thoughts**.

Input-Output Distribution Alignment

To mitigate the distribution shift between input token embedding and last-layer hidden state, a linear alignment operator is employed. A projection matrix $W_a \in \mathbb{R}^{d_h \times d_h}$ maps output vector $h \in H$ to a new input vector e .

$$e = h W_a, \quad \text{where } W_a \approx W_{\text{out}}^{-1} W_{\text{in}}$$

Latent Communication

Latent Working Memory

Let A_1, A_2 be two agents. A_1 first performs m latent steps and obtain KV Cache or latent memory \mathcal{M}_{A_1} .

$$\mathcal{M}_{A_1} = \left\{ \left(K_{A_1, \text{cache}}^{(l)}, V_{A_1, \text{cache}}^{(l)} \right) \mid l = 1, 2, \dots, L \right\}$$

Information Transfer

Layer-wise concatenate KV caches from \mathcal{M}_{A_1} to A_2 by prepending each $K_{A_1, \text{cache}}^{(l)}$ and $V_{A_1, \text{cache}}^{(l)}$ to $K_{A_2, \text{cache}}^{(l)}$ and $V_{A_2, \text{cache}}^{(l)}$.

Introduction

Method

Experiment

Discussion

Multi-agent Setting

Sequential

Chain-of-agents design:
planer, critic, refiner, solver
(judge).

Hierarchical

Domain-specialized design:
code, math, science agents to
a summarizer.

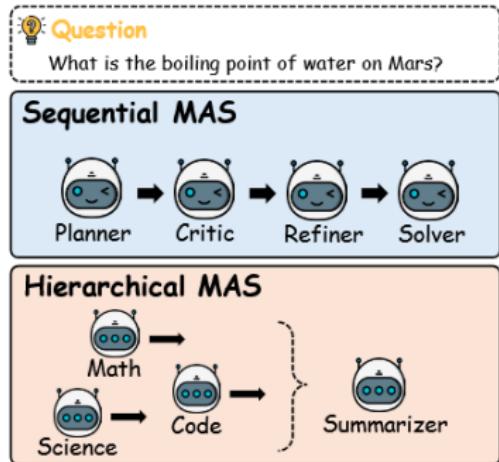


Figure 2 | Illustration of sequential and hierarchical MAS.

Results

Table 1 | Main results of LatentMAS on 6 general tasks under the Sequential MAS setting. We report 3 metrics in total, including task accuracy (%，“Acc.”), total output token usage (“Token”), and end-to-end inference speed (time(s) / run, “Speed”). We compare LatentMAS with both TextMAS and single-model (“Single”) baselines. For each metric, we **bold** the better performance and visualize LatentMAS gains over TextMAS in the **Improve** columns.

Tasks	Metrics	Qwen3-4B			Improve	Qwen3-8B			Improve	Qwen3-14B			Improve
		Single	TextMAS	LatentMAS		Single	TextMAS	LatentMAS		Single	TextMAS	LatentMAS	
<i>Sequential MAS Setting</i>													
ARC-E	Acc.	95.4	96.4	98.6	$\uparrow 2.2$	95.6	99.1	98.8	$\downarrow 0.3$	97.2	99.0	99.4	$\uparrow 0.4$
	Token	724	2420	581	$\downarrow 76.0\%$	656	2085	490	$\downarrow 76.5\%$	608	1670	224	$\downarrow 86.6\%$
	Speed	369	2874	512	$\times 5.6$	404	3702	1759	$\times 2.1$	551	9171	2124	$\times 4.3$
ARC-C	Acc.	89.2	90.0	92.3	$\uparrow 2.3$	91.0	94.6	94.4	$\downarrow 0.2$	92.6	95.9	95.6	$\downarrow 0.3$
	Token	913	2678	718	$\downarrow 73.2\%$	846	2252	529	$\downarrow 76.5\%$	773	2985	426	$\downarrow 85.7\%$
	Speed	97	1579	260	$\times 6.1$	266	2059	703	$\times 2.9$	338	5125	1136	$\times 4.5$
GSM8K	Acc.	82.4	89.8	88.2	$\downarrow 1.6$	81.1	92.3	93.8	$\uparrow 1.5$	83.7	93.8	95.2	$\uparrow 1.4$
	Token	1136	3172	607	$\downarrow 80.9\%$	1280	2324	860	$\downarrow 63.0\%$	1118	3324	644	$\downarrow 80.6\%$
	Speed	469	1970	375	$\times 5.3$	449	1739	543	$\times 3.2$	536	3729	1952	$\times 1.9$
MedQA	Acc.	47.7	65.3	66.3	$\uparrow 1.0$	53.0	75.0	75.3	$\uparrow 0.3$	64.7	80.3	80.7	$\uparrow 0.4$
	Token	2134	3962	1685	$\downarrow 57.5\%$	2098	4260	1555	$\downarrow 63.5\%$	1746	3444	1841	$\downarrow 46.5\%$
	Speed	236	1267	438	$\times 2.9$	476	1923	928	$\times 2.1$	1360	4142	1420	$\times 2.9$
MBPP+	Acc.	63.5	69.8	73.5	$\uparrow 3.7$	64.8	69.5	74.6	$\uparrow 5.1$	68.5	72.8	75.7	$\uparrow 2.9$
	Token	1634	4420	1339	$\downarrow 69.7\%$	2053	3695	1164	$\downarrow 68.5\%$	1858	4971	1621	$\downarrow 67.4\%$
	Speed	523	2148	577	$\times 3.7$	1064	3628	1275	$\times 2.8$	2410	8728	2400	$\times 3.6$
HumanEval+	Acc.	75.0	79.7	79.9	$\uparrow 0.2$	74.4	80.5	80.5	$\uparrow 0.0$	76.8	81.1	86.5	$\uparrow 5.4$
	Token	2380	5987	1775	$\downarrow 70.4\%$	2507	4593	1866	$\downarrow 59.4\%$	2366	5934	2042	$\downarrow 65.6\%$
	Speed	274	1044	350	$\times 3.0$	502	1619	497	$\times 3.3$	1084	4062	1285	$\times 3.2$

Table 2 | Main results of LatentMAS on 6 general tasks under the Hierarchical MAS setting. We report accuracy, token usage, and end-to-end speed, and highlight the performance gains following the same evaluation protocol as in Table 1.

Tasks	Metrics	Qwen3-4B			Improve	Qwen3-8B			Improve	Qwen3-14B			Improve
		Single	TextMAS	LatentMAS		Single	TextMAS	LatentMAS		Single	TextMAS	LatentMAS	
<i>Hierarchical MAS Setting</i>													
ARC-E	Acc.	95.4	97.1	96.8	↓ 0.3	95.6	98.2	98.3	↑ 0.1	97.2	98.3	98.7	↑ 0.4
	Token	724	2054	363	↓ 82.3%	656	2237	308	↓ 86.2%	608	2752	619	↓ 77.5%
	Speed	369	2239	591	×3.8	404	3619	1779	×2.0	551	7102	1884	×3.8
ARC-C	Acc.	89.2	92.5	91.7	↓ 0.8	91.0	93.3	93.9	↑ 0.6	92.6	95.3	95.5	↑ 0.2
	Token	913	2674	447	↓ 83.3%	846	2854	344	↓ 87.9%	773	2167	295	↓ 86.4%
	Speed	97	1275	299	×4.3	266	2034	714	×2.8	338	4283	1090	×3.9
GSM8K	Acc.	82.4	89.4	88.4	↓ 1.0	81.1	90.4	89.5	↓ 0.9	83.7	90.8	91.6	↑ 0.8
	Token	1136	3098	555	↓ 82.1%	1280	2370	353	↓ 85.1%	1118	3021	495	↓ 83.6%
	Speed	469	1878	360	×5.2	449	1365	702	×1.9	536	3675	1631	×2.3
MedQA	Acc.	47.7	65.0	67.3	↑ 2.3	53.0	76.3	77.0	↑ 0.7	64.7	78.0	78.3	↑ 0.3
	Token	2134	6702	1015	↓ 84.9%	2098	6893	1007	↓ 85.4%	1746	5473	899	↓ 83.6%
	Speed	236	1495	557	×2.7	476	3387	964	×3.5	1360	7591	1250	×6.1
MBPP+	Acc.	63.5	69.3	70.6	↑ 1.3	64.8	71.9	72.2	↑ 0.3	68.5	73.0	73.8	↑ 0.8
	Token	1634	6782	1339	↓ 80.3%	2053	7703	1264	↓ 83.6%	1858	7458	1187	↓ 84.1%
	Speed	523	1766	489	×3.6	1064	3898	1387	×2.8	2410	9162	2507	×3.7
HumanEval+	Acc.	75.0	76.2	79.3	↑ 3.1	74.4	76.8	78.0	↑ 1.2	76.8	84.1	86.6	↑ 2.5
	Token	2380	8127	1373	↓ 83.1%	2507	8768	1274	↓ 85.5%	2366	8114	1512	↓ 81.4%
	Speed	274	931	333	×2.8	502	1809	439	×4.1	1084	3988	1188	×3.4

Table 3 | Main results of LatentMAS on 3 reasoning-intensive tasks under both Sequential and Hierarchical MAS settings. We report accuracy, token usage, and end-to-end speed, and highlight the performance gains following the same evaluation protocol as in Table 1.

Tasks	Metrics	Qwen3-8B			Improve	Qwen3-14B			Improve
		Single	TextMAS	LatentMAS		Single	TextMAS	LatentMAS	
<i>Sequential MAS Setting</i>									
AIME24	Acc.	50.0	53.3	56.7	↑ 3.4	63.3	63.3	66.7	↑ 3.4
	Token	12891	38596	8953	↓ 76.8%	11263	32092	10593	↓ 67.0%
	Speed	421	2808	688	×4.1	1018	4554	1149	×4.0
AIME25	Acc.	46.7	53.3	53.3	↑ 0.0	56.7	60.0	63.3	↑ 3.3
	Token	14692	45088	8699	↓ 80.7%	11298	44618	11402	↓ 74.4%
	Speed	450	3150	820	×3.8	1040	5184	1473	×3.5
GPQA-Diamond	Acc.	39.9	43.4	45.5	↑ 2.1	48.5	51.5	52.0	↑ 0.5
	Token	6435	17986	4571	↓ 74.6%	5547	12676	5454	↓ 57.0%
	Speed	813	5771	854	×6.8	1043	9714	1475	×6.6
<i>Hierarchical MAS Setting</i>									
AIME24	Acc.	50.0	53.3	53.3	↑ 0.0	63.3	70.0	73.3	↑ 3.3
	Token	12891	42629	7526	↓ 82.3%	11263	29025	10230	↓ 64.8%
	Speed	421	3132	776	×4.0	1018	5718	1089	×5.3
AIME25	Acc.	46.7	50.0	50.0	↑ 0.0	56.7	66.7	66.7	↑ 0.0
	Token	14692	53929	13230	↓ 75.5%	11298	50003	9527	↓ 80.9%
	Speed	450	3488	616	×5.7	1040	6019	1056	×5.7
GPQA-Diamond	Acc.	39.9	43.0	46.9	↑ 3.9	48.5	52.0	53.0	↑ 1.0
	Token	6435	22450	3395	↓ 84.9%	5547	20931	3606	↓ 82.8%
	Speed	813	6108	798	×7.7	1043	9119	1458	×6.3

Study on Input-Output Alignment

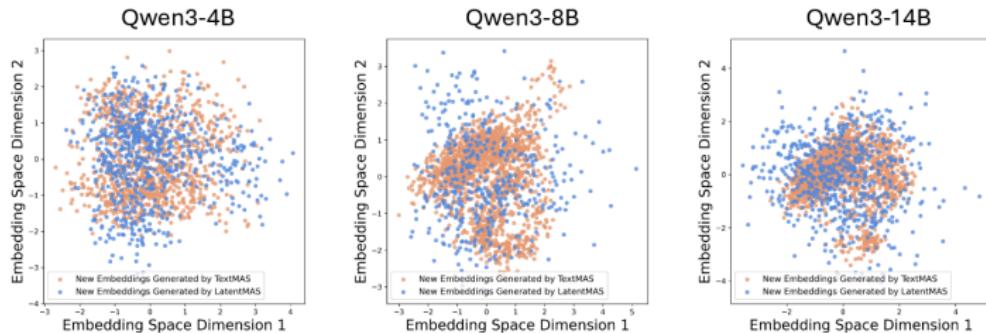


Figure 5 | Illustration of the semantic meaning encoded by latent thoughts in LatentMAS. Newly generated latent thought embeddings in LatentMAS largely cover the embedding space of text-based generated tokens, indicating semantic consistency and greater expressive capacity than discrete text.

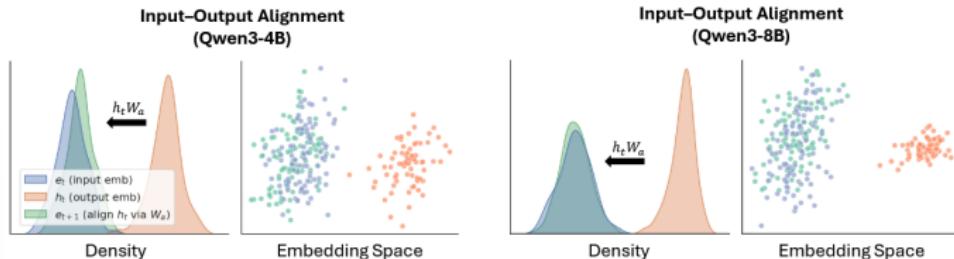


Figure 6 | Effectiveness of the input-output alignment W_a on MedQA. Unaligned output embeddings (h_t) drift away from the original input embeddings (e_t), while the aligned vectors (e_{t+1}) realign with e_t , demonstrating that W_a preserves embedding-space structure and prevents representation drift.

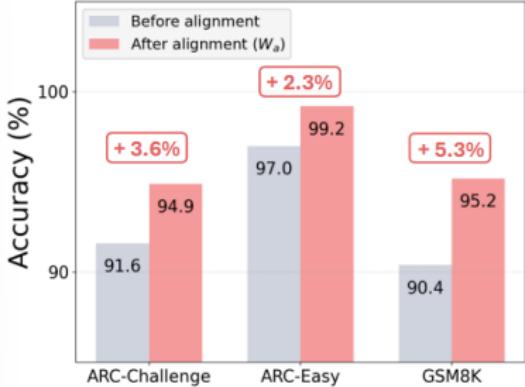


Figure 7 | Downstream performance before/after applying the input-output alignment W_a .

m	0	10	20	40	80	160
ARC-C	91.3	93.4	93.4	94.9	94.8	93.7
ARC-E	94.7	98.9	98.9	99.4	99.6	98.3
GSM8K	85.6	90.3	90.9	91.4	92.0	91.9

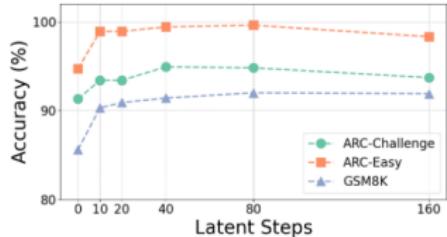


Figure 8 | Effectiveness of different latent step depths of LatentMAS on downstream performance.

Introduction

Method

Experiment

Discussion

Expressiveness of Latent Thoughts

Theorem 3.1 (Expressiveness of Latent Thoughts). *Under the Linear Representation Hypothesis on h (detailed in Assumption B.1), if the sequence of all latent thoughts with length m can be expressed losslessly through corresponding text-based reasoning, then the length of text (in tokens) needs to be at least $\Omega(d_h m / \log |\mathcal{V}|)$, where $|\mathcal{V}| > 1$ denotes the vocabulary size.*

Remark 3.2. Theorem 3.1 suggests that latent thoughts generation can be $O(d_h / \log |\mathcal{V}|)$ times more efficient than text-based reasoning. In addition, the expressiveness scales linearly with d_h , implying that larger models inherently exhibit greater latent reasoning capacity.

- ▶ This theorem suggests that, if the latent thoughts can be limited to a fixed set (combination of semantic basis), then to represent all latent thoughts with text token embeddings in a limited vocabulary size, there should exist a combination method.
- ▶ This theorem is only meaningful when we assume that multiple text token embeddings can be equivalent to a latent embedding during generation.

Reproduction Issue

The results in the table can not be fully reproduced. Here is my results on AIME 2024 under sequential setting:

- ▶ Single agent: 76.7%
- ▶ Text-MAS: 43.3%
- ▶ Latent-MAS (10 step): 66.7%
- ▶ Latent-MAS (20 step): 46.7%

Analysis

- ▶ Due to the uncertainty in decoding, latent reasoning is more deterministic than text token.
- ▶ The multi-agent setting is not reasonable, which causes lower performance than baseline.
- ▶ The latent steps, as a hyperparameter, have great impacts on the final performance.

Conclusion

Contribution

- ▶ Build a framework that combines latent reasoning and latent communication.
- ▶ Propose the input-output alignment method.

Problems

- ▶ The reproduction issue should be investigated.
- ▶ As an agent system, latent reasoning can not coexist with tool using. The workflow is not dynamic.